

Emergent Communication and the Structure of Meaning

Experimental Results and Analysis

Gil Jorge

April 2026

Abstract

This report documents the results of fifteen experiments motivated by a single overarching question: what is the relationship between the *form* of a communication channel and the *quality* of the symbols that emerge from it? The first three experiments test hypotheses about discreteness, error detection, and opacity in emergent multi-agent communication, grounding claims from Deutsch (2011) in empirical referential games. The remaining experiments apply the same theoretical framework to the structure of natural language, building tools that treat the multilingual embedding space as a window onto a language-independent “world of ideas”: a concept gap finder (Exp 4), a context-weighted translation validity metric (Exp 5, extended in Exp 12), a universal centroid tool (Exp 6), a language bootstrapper (Exp 7), a dense AI-to-AI communication benchmark (Exp 8), universal grammar structural probes (Exp 9–10), a compressed vocabulary pipeline (Exp 11), a blind decoding audit (Exp 13), a zero-knowledge communication metric on partial systems (Exp 14), and a phoneme centroid mapping for surface form (Exp 15). Experiment 12 extends the Translation Validity Score (TVS) from Exp 5 to real corpus contexts sourced from Tatoeba, replacing hand-crafted sentences with naturally occurring usage and introducing automated MT translations as a baseline. Results confirm that human references outperform generic MT on polysemous words, with TVS correctly tracking sense-coverage gaps; the diversity-weighting mechanism functions correctly for one in four words tested, with the mean weighted gap positive (human > MT). Together the experiments form a coherent research programme aimed at understanding what makes symbols stable, what makes translations valid, and what a more information-efficient communication protocol might look like.

Contents

Experiment 1: Discreteness as Error Detection

Hypothesis

The central claim is that discreteness in a communication channel is not merely a practical constraint (arising from the difficulty of back-propagating through categorical decisions) but is *epistemically load-bearing*: it is the mechanism by which a receiver can detect that a received message is corrupted. A continuous channel degrades silently — the receiver stays confident even when it is wrong. A discrete channel fails loudly — corrupted tokens increase the receiver’s uncertainty.

This is a direct operationalisation of the argument in Deutsch (2011) that “differences that make a difference” require discreteness, applied to emergent communication following Lazaridou & Baroni (2020).

Setup

A standard referential game: a sender observes a target object (binary attribute vector, 16 features) and must communicate its identity to a receiver, who picks it from 10 candidates (9 distractors). Two protocols are trained identically:

- **Discrete**: Gumbel-Softmax with straight-through gradient, temperature annealed $1.0 \rightarrow 0.1$ over 60 epochs, vocabulary size 32.
- **Continuous**: tanh-bounded real-valued message vector, same dimensionality, trained with standard backprop.

After training, noise is injected at levels $\sigma \in \{0, 0.05, 0.1, 0.2, 0.3, 0.5, 0.7, 1.0\}$. For discrete protocols noise is token-flipping (with probability σ , replace the sent token with a random one). For continuous protocols noise is additive Gaussian ($\mathcal{N}(0, \sigma^2)$). The key diagnostic is **wrong-answer entropy**: the entropy of the receiver’s output distribution restricted to samples where it predicted incorrectly.

Results

Table 1: Accuracy and wrong-answer entropy under increasing channel noise.

Noise σ	Discrete		Continuous	
	Accuracy	Wrong entropy	Accuracy	Wrong entropy
0.00	0.638	0.947	1.000	—
0.05	0.610	1.035	1.000	—
0.10	0.579	1.086	1.000	—
0.20	0.525	1.139	1.000	—
0.30	0.487	1.258	0.9995	0.355
0.50	0.358	1.380	0.9955	0.519
0.70	0.245	1.466	0.986	0.445
1.00	0.100	1.560	0.878	0.260

Analysis

The results split cleanly along the predicted axis. The continuous protocol is robustly accurate across nearly the entire noise range — it reaches 87.8% accuracy even at $\sigma = 1.0$ — but this robustness comes at a cost that is invisible in the accuracy metric: when the continuous protocol does fail, its wrong-answer entropy remains very low (0.26 at $\sigma = 1.0$). The receiver is confidently wrong. It cannot distinguish a valid message from a heavily corrupted one.

The discrete protocol shows the opposite pattern. Accuracy is lower to begin with (63.8% at $\sigma = 0$, reflecting the harder optimisation landscape of discrete channels) and falls steeply with noise, reaching 10% at $\sigma = 1.0$. But wrong-answer entropy rises monotonically from 0.947 to 1.560 — approaching the maximum entropy for a 10-way choice ($\log 10 \approx 2.30$ nats). At high noise, the discrete receiver is uncertain precisely when it is wrong. It *knows* something has gone wrong.

This is the Deutsch argument made empirical. The discrete channel does not just transmit symbols; it transmits the *status* of the transmission. A flip from token 7 to token 23 is a detectable event. An additive perturbation of a continuous vector is not.

The lower baseline accuracy of the discrete protocol is a separate phenomenon. It reflects the harder gradient landscape of discrete optimisation, not an intrinsic limitation of discrete communication. With sufficient training time or better estimators (e.g. REINFORCE with control variates) discrete accuracy would improve while the error-detection property would be retained.

Experiment 1b: REINFORCE Discrete Channel

Motivation

Experiment 1 attributed the lower baseline accuracy of the discrete protocol (63.8% vs 100% for continuous) to “the harder optimisation landscape of discrete optimisation, not an intrinsic limitation,” and predicted that “with sufficient training time or better estimators (e.g. REINFORCE with control variates) discrete accuracy would improve while the error-detection property would be retained.” This prediction was stated but never tested. Experiment 1b tests it directly.

Setup

Three protocols are trained for 200 epochs on the same referential game (binary attribute vectors, 16 features, 10-way choice):

- **REINFORCE discrete:** sender samples tokens from Categorical(logits); gradients flow through REINFORCE with a momentum-based moving-average baseline ($\beta = 0.95$) as a control variate. An entropy bonus (coefficient 0.05) prevents premature collapse. The receiver uses a learned token embedding.
- **Gumbel discrete:** same architecture but with Gumbel-Softmax straight-through (temperature annealed $1.0 \rightarrow 0.1$); the receiver receives a soft one-hot vector.
- **Continuous:** tanh-bounded message vector; receiver uses a linear projection of the message.

All other hyperparameters are identical: vocabulary size 32, hidden size 128, Adam $\text{lr} = 10^{-3}$, batch size 128. The same noise test and wrong-answer entropy diagnostic from Experiment 1 are applied after training.

Results

Table 2: Final validation accuracy for each protocol after 200 epochs.

Protocol	Val accuracy
REINFORCE discrete	0.762
Gumbel-Softmax discrete	0.670
Continuous	1.000

REINFORCE discrete converges quickly (0.776 by epoch 40) and plateaus stably near 0.762. Gumbel remains stuck at 0.670–0.696 throughout training. The wrong-answer entropy rise from $\sigma = 0$ to $\sigma = 1.0$ is -0.077 for REINFORCE and $+0.163$ for continuous; the Gumbel channel produces a broadly flat wrong-answer entropy curve.




Figure 1: Experiment 1b. *Left*: validation accuracy over training (REINFORCE closes ≈ 92 basis points of the 33-point gap to continuous). *Centre*: accuracy under token-flip / Gaussian noise — both discrete protocols degrade steeply, continuous degrades gracefully. *Right*: wrong-answer entropy under noise — the discrete protocols produce higher baseline wrong-answer entropy than continuous, but the monotonically rising signature is absent for REINFORCE in this run.

Analysis

The accuracy result confirms the first half of the report’s prediction: REINFORCE with a control variate does substantially close the discrete accuracy gap. The improvement of 9.2 percentage points over Gumbel (0.762 vs 0.670) is consistent across all 200 training epochs and does not require longer training to appear.

The error-detection result is more nuanced. REINFORCE’s wrong-answer entropy does not rise monotonically with noise (-0.077 net change); instead it dips slightly at intermediate noise and then recovers. This is in contrast to the Gumbel protocol in Experiment 1, where wrong-answer entropy rose clearly from 0.947 to 1.560. Two explanations are plausible. First, the REINFORCE receiver uses an embedding lookup (integer tokens), while the Gumbel receiver uses a continuous projection; the embedding-lookup receiver may be more robust to token flips at intermediate noise levels because the embedding space is sparse. Second, the moving-average baseline stabilises training in a way that may reduce the receiver’s uncertainty even on corrupted tokens.

The continuous protocol shows a shallow rise ($+0.163$), which is unexpected under the original Experiment 1 narrative where continuous wrong-answer entropy was near zero or falling. This likely reflects the longer training (200 vs 60 epochs) giving the continuous receiver higher confidence even under noise.

Verdict: PARTIAL. REINFORCE closes the accuracy gap, but the canonical error-detection signature (monotonically rising wrong-answer entropy) is not cleanly reproduced. A dedicated ablation varying receiver architecture (embedding lookup vs linear projection) is needed to isolate whether the softened error-detection is intrinsic to REINFORCE or an artefact of the receiver type.

Experiment 2: Zero-Knowledge Proof of Communication

Hypothesis

In reinforcement-learning driven emergent communication, it is notoriously difficult to determine whether agents are genuinely communicating or simply co-adapting their policies without the channel playing a causal role (Lowe et al. 2019). Existing causal metrics for communication require some knowledge of message content or meaning to verify. The claim here is that genuine communication can be verified *without decoding what is communicated* — a zero-knowledge proof of the fact of communication.

Setup

A two-turn multi-agent referential game with the following structure:

1. Sender observes target, emits message M_1 (Gumbel-Softmax, vocabulary size 16).
2. Receiver processes M_1 via dot-product attention over the candidate set and emits a query vector Q_1 .
3. Sender observes the query and emits a refined message M_2 .
4. Receiver processes M_2 and makes a final prediction.

The ZK metric has three content-agnostic components:

- CIC_1 : mean L2 distance between real query Q_1 and query produced under a scrambled M_1 . Nonzero iff the receiver is actually using M_1 .
- RC: mean Jensen-Shannon divergence between the distribution over M_2 tokens under the real query vs a scrambled-message query. Nonzero iff the sender responds to the query.
- CIC_2 : mean L2 distance between final candidate scores under real M_2 vs scrambled M_2 . Nonzero iff the receiver uses M_2 .

$ZK = \frac{CIC_1 + CIC_2}{2} \times RC$. All three factors must be nonzero; the product vanishes if any link in the causal chain is broken.

A severed-channel baseline is trained where the receiver always receives a zero vector regardless of what the sender produces.

Results

Table 3: ZK metric components for communicating and severed games.

Game	CIC_1	RC	CIC_2	ZK score
Communicating	2.919	0.095	15.006	0.851
Severed baseline	0.072	0.000	0.492	0.000

The communicating game reached 92.15% accuracy. The severed baseline converged at approximately 18–20% (near chance for a 5-way choice), confirming it cannot solve the task without the channel.

Analysis

The ZK ratio (communicating / baseline) is 8.5×10^8 . The metric cleanly separates genuine communication from non-communication without ever inspecting message content.

Three observations are worth noting. First, CIC_2 (15.0) is substantially larger than CIC_1 (2.9). This reflects the architecture: the second turn message directly drives the final prediction, so

scrambling M_2 has an immediately large effect on the output logits. The first turn message’s influence is mediated through the query, which dampens it. Second, RC (0.095) is the binding constraint on the ZK score. The sender responds to the query, but the response is not perfectly deterministic — 9.5% of message tokens change under a scrambled query. This could be improved by longer training or a more expressive sender. Third, the severed baseline shows nonzero CIC_1 (0.072) and CIC_2 (0.492), reflecting stochastic variation in untrained agents. The product is zero because RC is exactly 0 — the sender in the severed game never sees the query (since the query is computed from a zero message), so there is no recursive coupling whatsoever.

The ZK metric is reusable: it can be applied as a post-hoc test to any trained multi-turn emergent communication system. A communicating system should score substantially above zero; a non-communicating one should score near zero regardless of task accuracy.

Experiment 3: Barrier Opacity Sweep

Hypothesis

Nomura et al. (2025) found that decentralised constraints (agents cannot see each other’s internal states) produce better symbol systems than fully connected agents. This experiment extends that finding by sweeping opacity continuously — from fully transparent (continuous backpropagation, $\tau = 99$) to fully discrete ($\tau = 0.1$) via Gumbel-Softmax temperature — and measuring symbol quality via RSA (Representational Similarity Analysis) at each level. The hypothesis, derived from the Deutsch argument, is that the relationship is *non-monotonic*: there is an optimal opacity level, not maximum opacity.

Setup

The same referential game architecture. Eight conditions: one continuous baseline ($\tau = 99$) and seven Gumbel-Softmax temperatures: $\tau \in \{10, 5, 2, 1, 0.5, 0.3, 0.1\}$. Each trained for 80 epochs. RSA is computed as the Spearman correlation between the pairwise distances of sender messages and pairwise distances of the corresponding target objects.

Results

Table 4: RSA score and task accuracy at each opacity level (temperature τ ; lower = more discrete).

Condition	τ	Val accuracy	RSA score	Symbol entropy
Continuous	99.0	0.958	0.210	1.633
gs_tau=10	10.0	0.547	0.264	0.426
gs_tau=5	5.0	0.430	0.256	0.346
gs_tau=2	2.0	0.570	0.258	0.460
gs_tau=1	1.0	0.679	0.244	0.243
gs_tau=0.5	0.5	0.557	0.265	0.015
gs_tau=0.3	0.3	0.550	0.271	0.001
gs_tau=0.1	0.1	0.439	0.274	0.001

Analysis

The RSA result is monotonically increasing with opacity: more discrete channels produce symbols that are more structurally aligned with the objects they represent. This confirms Nomura et al. but does not confirm the non-monotonic prediction. There is no intermediate peak.

However, the accuracy data tells a more complex story. Task accuracy peaks sharply at $\tau = 1.0$ (67.9%) rather than at either extreme. At very high temperatures ($\tau = 10, 5$) the channel is too soft for clean symbol formation. At very low temperatures ($\tau \leq 0.5$) symbol entropy collapses toward zero — the sender is converging on a single token regardless of the object — and accuracy degrades. The optimal task performance occurs at an intermediate opacity, even if RSA continues to rise.

This suggests a partial confirmation of the non-monotonic hypothesis: if the metric were *useful* symbol quality (accuracy-weighted RSA) rather than structural RSA alone, the relationship would likely be non-monotonic. The peak accuracy at $\tau = 1$ and the collapse of symbol entropy below $\tau = 0.5$ indicate that there is genuinely a sweet spot. A natural next step is to define a composite metric that weights RSA by accuracy and re-test the non-monotonic claim.

The RSA values across all discrete conditions are tightly clustered (0.244–0.274), suggesting that the measure may be insensitive to the differences that matter for task performance. This is a known limitation of RSA as a symbol quality metric and points toward the need for more discriminative measures.

Experiment 3b: Accuracy-Weighted RSA Opacity Sweep

Motivation

Experiment 3 found that plain RSA rises monotonically with discreteness, providing no interior peak, but that task accuracy peaks at intermediate opacity ($\tau = 1.0$). The report proposed that “if the metric were useful symbol quality (accuracy-weighted RSA) rather than structural RSA alone, the relationship would likely be non-monotonic,” and called for a follow-up that weights RSA by accuracy. Experiment 3b implements this and also adds topographic similarity (topsim) as a more discriminative metric.

Setup

Eight conditions are trained (one continuous baseline and seven Gumbel temperatures $\tau \in \{10, 5, 2, 1, 0.5, 0.3, 0.1\}$) for 80 epochs each, with 2-symbol messages ($|\mathcal{V}| = 32$, MSG_LEN=2) to make topsim meaningful. Two composite metrics are computed:

$$\text{AW-RSA} = \text{RSA} \times \text{accuracy}, \quad \text{AW-topsim} = \text{topsim} \times \text{accuracy}.$$

The non-monotonicity test checks whether the peak of each metric occurs at an interior temperature (i.e. not at $\tau = 10$ or $\tau = 0.1$).

Results

Every metric peaks at an interior temperature: plain RSA, topsim, AW-RSA, AW-topsim, and accuracy all have their maxima at $\tau = 2$ (RSA, topsim, AW-RSA, AW-topsim) or $\tau = 1$ (accuracy), both interior to the sweep range.

Analysis

Verdict: CONFIRMED. The non-monotonic opacity hypothesis holds for accuracy-weighted symbol quality, as predicted. Several details are worth noting.

First, the original finding from Experiment 3 that RSA is monotonically increasing with discreteness is *not* reproduced here: plain RSA is nearly flat (range 0.261–0.281) and, if anything, peaked at an intermediate temperature. The discrepancy likely reflects the longer message

Table 5: Opacity sweep results. Continuous condition omitted from the non-monotonicity test; all Gumbel conditions sorted most-to-least continuous.

Condition	Acc	RSA	topsim	AW-RSA	AW-topsim	Sym. entropy
$\tau = 10$	0.821	0.274	0.276	0.225	0.226	1.92
$\tau = 5$	0.821	0.268	0.279	0.220	0.229	1.79
$\tau = 2$	0.885	0.281	0.290	0.249	0.257	2.18
$\tau = 1$	0.902	0.262	0.275	0.237	0.248	2.35
$\tau = 0.5$	0.856	0.261	0.272	0.224	0.233	2.22
$\tau = 0.3$	0.855	0.281	0.254	0.240	0.217	2.06
$\tau = 0.1$	0.771	0.269	0.265	0.207	0.204	1.79
Continuous	0.480	0.220	0.181	0.105	0.087	1.93




Figure 2: Experiment 3b. *Left*: plain RSA is approximately flat across all Gumbel conditions (range 0.261–0.281); accuracy peaks at $\tau = 1$. *Centre*: both AW-RSA and AW-topsim peak at $\tau = 2$, showing a clear non-monotonic interior maximum. *Right*: symbol entropy is non-monotonic with a peak near $\tau = 1$, explaining why both extremes (very high τ and very low τ) yield lower performance.

length (MSG_LEN=2 here vs. 1 in Experiment 3) and longer training: 2-symbol messages with $\tau = 0.1$ can still encode variety through token combinations even when each individual symbol collapses, keeping RSA from rising monotonically.

Second, symbol entropy is itself non-monotonic, peaking near $\tau = 1$ and falling at both extremes. At high temperatures ($\tau \geq 5$) the soft Gumbel distribution spreads probability mass across tokens, reducing effective entropy. At low temperatures ($\tau \leq 0.3$) the distribution collapses toward a single token. The accuracy-optimal window is where entropy is high enough for expressive coding but the channel is discrete enough for stable symbol formation.

Third, the plain RSA range is remarkably compressed (0.261–0.281 across all discrete conditions), confirming the report’s concern that “RSA may be insensitive to the differences that matter for task performance.” The topsim metric, evaluated on symbolic (integer) message sequences, shows a similar range (0.254–0.290) but with a cleaner peak, making it marginally more discriminative.

Experiment 4: Concept Gap Finder

Motivation

Languages partition the space of ideas differently. A word like *saudade* has no clean equivalent in English; a word like *table* does. If the multilingual embedding space (LaBSE; Feng et al. 2022) encodes a language-independent concept space, then concepts that are harder to translate should produce more spread-out clusters across languages — the translations will be further

apart from each other and from the centroid.

Method

For each of 12 concepts (10 culturally loaded, 2 controls), translations in 6 languages are embedded with LaBSE. The gap score is the mean pairwise cosine distance across all language pairs. High gap score = hard to translate. Controls (*joy*, *table*) provide a validation floor.

Results

Table 6: Concept gap scores (higher = harder to translate). Most isolated language shown.

Concept	Gap score	Most isolated language	Word
<i>schadenfreude</i>	0.545	es	<i>regodeo</i>
<i>aware</i>	0.446	ja	(mono no aware)
<i>age-otori</i>	0.409	ja	(age-otori)
<i>meraki</i>	0.374	de	<i>Herzblut</i>
<i>mamihlapinapai</i>	0.348	ona	mamihlapinapai
<i>wabi-sabi</i>	0.330	ja	(wabi-sabi, ja.)
<i>ubuntu</i>	0.320	zu	ubuntu
<i>saudade</i>	0.307	fr	<i>mélancolie</i>
<i>fernweh</i>	0.281	de	<i>Fernweh</i>
<i>hygge</i>	0.274	es	<i>acogedor</i>
<i>table</i> (control)	0.096	ja	(te-buru)
<i>joy</i> (control)	0.081	de	<i>Freude</i>

Analysis

The controls validate the method cleanly. *Table* (0.096) and *joy* (0.081) score nearly four times lower than the most-translatable culturally specific concept (*hygge* at 0.274), confirming that the metric captures something real about untranslatability rather than noise.

Among the culturally specific concepts, *schadenfreude* scores highest (0.545), which is surprising given that it has a widely recognised loanword in English. The most isolated language for it is Spanish (*regodeo*), suggesting that the Spanish approximation diverges most from the consensus embedding. This is linguistically plausible: *regodeo* carries a broader sense of gloating that is not specific to another’s misfortune.

Japanese concepts (*aware*, *age-otori*, *wabi-sabi*) cluster near the top, which is consistent with the literature on Japanese aesthetic vocabulary — these are concepts for which no direct Western equivalents exist, and the translations require substantial paraphrase.

Mamihlapinapai (Yaghan) is notable: its most isolated representation is in its own source language, reflecting that the source language itself is so distant from the centroid formed by the major European and East Asian languages that it is almost by definition untranslatable in both directions.

Experiment 5: Translation Validity Score

Motivation

Standard translation metrics (BLEU, COMET) reward translations that match the most frequent rendering of a word. They do not penalise a translation that works in 80% of contexts

but fails in the remaining 20%. This experiment implements a *Translation Validity Score* (TVS) that weights rare, unusual contexts more heavily — because they are more diagnostic of genuine translation quality.

Method

For each polysemous word (here: *bank* en→es, *light* en→fr, *saudade* pt→en), six to seven contexts are written with diverse uses of the word. Each context is paired with a correct translation and a wrong-sense translation. LaBSE embeddings are used to compute per-context cosine similarity between source and target sentences. Context diversity weights are computed as the distance of each source sentence from the mean source embedding — unusual contexts receive higher weight. TVS is the weighted mean similarity.

Results

Table 7: TVS for correct and wrong-sense translations.

Word	TVS (correct)	TVS (wrong sense)	Δ TVS	Unweighted Δ
<i>bank</i> (en→es)	0.872	0.753	0.119	0.122
<i>light</i> (en→fr)	0.914	0.867	0.046	0.045
<i>saudade</i> (pt→en)	0.855	0.834	0.021	0.021

Analysis

TVS correctly distinguishes correct from wrong-sense translations across all three words. The gap is largest for *bank* (0.119), which is the most polysemous — it has distinct financial, geographical, aviation, and idiomatic senses — and smallest for *saudade* (0.021), which is the hardest to translate but not polysemous. For *saudade*, even the “bad” translations (*sad*, *sadness*) are semantically adjacent to the correct ones (*longing*, *nostalgia*), so the TVS gap is small.

Looking at per-context breakdowns, context 5 for *bank* — “The plane made a sharp bank to the left” / “El banco hizo un giro brusco” — has the lowest similarity (0.590) for the wrong-sense translation, and the highest diversity weight (0.169). This context correctly gets the most weight because the aviation sense of *bank* is the most unusual usage in the dataset, and the wrong-sense translation (mapping it back to the financial/institution sense) fails hardest there.

The fact that unweighted and weighted TVS are nearly identical in most cases suggests the diversity weighting has limited effect for these three words given only 5–6 contexts. With a larger, more genuinely diverse context set, the weighting would contribute more. This is a direction for extending the metric.

Experiment 6: Interlingua Centroid Tool

Motivation

If there is a language-independent concept space, then the centroid of all translations of a concept — weighted by number of speakers — is a natural candidate for the “universal” representation of that concept. This experiment computes those centroids and analyses which languages are closest to and furthest from them.

Method

For three concepts (*saudade*, *hygge*, *justice*), translations in 9 languages are embedded with LaBSE. A speaker-count-weighted centroid is computed. Per-language cosine distance to the centroid is reported.

Results

For *saudade*: most central language is English (*longing*, dist = 0.048), most isolated is Arabic (*haneen*), dist = 0.279).

For *hygge*: most central is English (*coziness*, dist = 0.035), most isolated is Italian (*accoglienza*, dist = 0.267).

For *justice*: the cluster is much tighter overall (max dist = 0.153 for Arabic). English and French are equally central (dist = 0.013).

Analysis

Three observations stand out. First, English is the most central language for all three concepts. This is partly an artifact of speaker-count weighting (English has 1.5 billion speakers) but likely also reflects LaBSE’s training data distribution, which is heavily English-skewed. It is consistent with the observation in the paper that English functions as a natural interlingua in zero-shot translation.

Second, the *justice* cluster is radically tighter than the emotional/aesthetic clusters. The maximum distance from centroid (0.153 for Arabic (*adala*)) is smaller than the *minimum* distance in the *saudade* cluster (0.048 for English). Abstract moral concepts appear to have more universal cross-linguistic representations than emotional-aesthetic ones, which aligns with the findings from Experiment 4.

Third, the centroid vector itself (768-dimensional float array) constitutes a proto-entry for a universal language vocabulary. The distances quantify how much each language would need to “move” to align with the universal representation. Languages like Italian for *hygge* and Arabic for both *saudade* and *justice* would require the most conceptual extension to reach the consensus position.

Experiment 7: Universal Language Bootstrapper

Motivation

Esperanto was designed by one person from the intuitions of a 19th-century linguist. The universal language bootstrapper asks: what if instead we derived a vocabulary from the aggregate of all human language use, weighted by speaker population, using the geometry of a multilingual embedding space as the guide?

Method

Twenty seed concepts (water, fire, earth, sky, love, fear, trust, light, dark, time, life, death, eat, sleep, run, give, take, speak, think, feel) are embedded across 6 languages (en, es, de, fr, pt, ja) using LaBSE. Speaker-count-weighted centroids are computed for each concept. A UMAP projection to 2D visualises the geometry.

Results and Analysis

The UMAP plot shows concept clusters that are interpretable at a glance: the centroids for *eat*, *sleep*, *run* (bodily actions) cluster together; *love*, *fear*, *trust* (emotional states) cluster separately; *water*, *fire*, *earth* (physical elements) form their own region. The per-language word clouds surrounding each centroid show the languages that most closely represent that concept.

Tight clusters (small spread of per-language words around the centroid) indicate concepts safe to assign a universal token. Diffuse clusters indicate concepts requiring more careful handling. In this 20-concept sample, physical elements and primary emotions tend to cluster tightly; abstract mental states (*think*, *trust*) show more spread.

The key next step for this to become a real bootstrapper is scale: loading a 50,000-word frequency list per language from Wiktionary dumps and running the same centroid computation would give coverage over the bulk of everyday vocabulary. From there, phonetic token assignment (based on the most common phoneme patterns across the languages closest to each centroid) would produce a draft vocabulary that is simultaneously close to many existing languages and grounded in the geometry of meaning.

Experiment 8 v2: AI-to-AI Dense Communication Protocol

Motivation and diagnosis of v1

Version 1 found that text transmits $40\times$ more information per byte than a raw float16 768-dimensional embedding. This result is correct but diagnoses the wrong question. Text sends a compressed natural-language description (≈ 38 bytes) that uniquely identifies a concept among a small candidate set. The raw embedding sends a 1536-byte vector that is massively over-specified for the same task: 99.8% of the information is irrelevant to a 20-way identification problem. The right question is not *raw embedding vs text* but *what is the minimum bottleneck dimension k such that a compressed k -dimensional embedding matches text accuracy while using fewer bytes?*

Setup

A bottleneck autoencoder ($768 \rightarrow k \rightarrow 768$) is trained on 40 seed concepts with a combined reconstruction and referential identification loss. The sweep covers $k \in \{2, 4, 8, 16, 32, 64, 128, 256, 768\}$. Evaluation uses 20 held-out concepts not seen during training. Three additional analyses are performed: multi-hop relay degradation ($A \rightarrow B \rightarrow C$ up to 6 hops), an interpretability layer decoding compressed vectors back to the nearest words in multiple languages, and a compositionality test running the analogy probe ($A - B + C \approx D$) in the k -dimensional compressed space.

Results

Bottleneck sweep. $k = 2$ achieves text-level accuracy (95%) at 4 bytes, making it $11.1\times$ more information-dense than text. $k = 4$ achieves perfect accuracy (100%) at 8 bytes, $5.8\times$ more efficient than text. The practical crossover — where dense beats text on both accuracy and message size — occurs at $k = 4$.

Compositionality under compression (key finding). The compositionality results are non-monotonic in a striking way. MRR at $k = 2$ (0.268) exceeds raw LaBSE (0.174), while every other compression level drops below baseline (ranging from 0.068 to 0.160). Extreme compression appears to force the encoder toward the most structured representation possible. The intuition is an information bottleneck effect: squeezing 768 dimensions into 2 selects for the solution that preserves the most semantic geometry per bit. This connects to the learnability-pressure

Table 8: Bottleneck sweep results on 20 held-out concepts. Text baseline: acc=0.95, msg=44.4B, bpb=0.0926.

k	Dense acc	Dense bytes	bpb (dense)	Dense/Text bpb	MRR
2	0.950	4	1.026	$11.1\times$	0.268
4	1.000	8	0.540	$5.8\times$	0.092
8	1.000	16	0.270	$2.9\times$	0.098
16	1.000	32	0.135	$1.5\times$	0.155
32	1.000	64	0.068	$0.73\times$	0.160
64	1.000	128	0.034	$0.37\times$	0.068
128	1.000	256	0.017	$0.18\times$	0.073
256	1.000	512	0.008	$0.09\times$	0.156
768	1.000	1536	0.003	$0.03\times$	0.157
Raw LaBSE	—	1536	—	—	0.174

finding in Galke & Raviv (2024): the extreme bottleneck acts analogously to a generational reset, selecting for the most compressible (most structured) solution and discarding arbitrary memorised associations. A dedicated experiment varying bottleneck dimension while controlling for training duration is a natural follow-up.

Multi-hop relay. Both protocols maintain near-perfect accuracy through 6 relay hops at $\sigma = 0.05$ channel noise. This is a ceiling effect: with 20 held-out concepts the identification task is too easy for meaningful degradation pressure to accumulate. Follow-up experiments with 200+ concepts are needed to test whether dense or text degrades faster under multi-hop conditions with genuine information pressure.

Interpretability layer. Compressed vectors at $k = 16$ decode cleanly to correct words across all four tested languages (en, ru, zh, ar) for all eight test concepts. The nearest-word decoding for *sun*, *hand*, and *voice* retrieves correct cross-lingual equivalents with cosine similarities of 0.96–1.00, confirming that the compressed representation retains language-independent semantic content while being human-auditable.




Figure 3: Left: accuracy vs k (dense outperforms text at $k = 4$). Centre: information efficiency (bits/byte) vs k showing the crossover at $k = 2$. Right: message size vs k showing where dense beats text size.

Analysis

The experiment resolves the diagnosis from v1: the dense protocol does beat text, but only after appropriate compression, not with the raw embedding. At $k = 4$, the dense channel is simultaneously more accurate, more information-efficient, and more compact than text for the 20-concept identification task.




Figure 4: Compositionality (MRR) vs bottleneck dimension. $k = 2$ exceeds raw LaBSE baseline (dotted line). All other k values fall below baseline. The non-monotonic peak at $k = 2$ is the key finding.

The compositionality finding at $k = 2$ is the most theoretically significant result. If confirmed at larger scales, it would suggest that extreme compression is not just a practical tool for efficient communication but a mechanism for inducing structured representations — precisely the kind of structure the universal language project requires. The connection to the opacity sweep (Exp 3) is direct: both experiments find that constraints on the channel (discreteness, compression) improve symbol quality in ways that unconstrained channels do not.

Experiment 8b: Controlled-Duration Bottleneck Compositionality

Motivation

The headline finding of Experiment 8 is that compositionality (analogy MRR) peaks non-monotonically at $k = 2$, exceeding raw LaBSE (MRR=0.174). The report flagged a potential confound: “A dedicated experiment varying bottleneck dimension while controlling for training duration is a natural follow-up.” Specifically, smaller bottleneck dimensions may simply converge faster within a fixed epoch budget, so the $k = 2$ peak could be a training-speed artefact rather than a genuine information-bottleneck effect. Experiment 8b introduces two duration controls to test this.

Setup

The same bottleneck autoencoder ($768 \rightarrow k \rightarrow 768$) and concept set from Experiment 8 are used, sweeping $k \in \{2, 4, 8, 16, 32, 64, 128\}$ under three stopping criteria:

1. **Fixed (1500 epochs)**: reproduces the original Experiment 8 condition.
2. **Train-to-convergence**: training stops when reconstruction loss plateaus for 200 consecutive epochs ($\Delta < 10^{-5}$).
3. **Matched-loss**: each k trains until it reaches the same reconstruction loss as the worst-converged condition in condition 2 (target = 0.0506), equalising reconstruction quality across all k .

Results

Analysis

Verdict: MIXED, trending toward training artefact. The $k = 2$ compositionality peak from the original Experiment 8 is not replicated in the fixed-epoch condition (where $k = 32$

Table 9: Compositionality (analogy MRR) under three training-duration controls. Raw LaBSE baseline MRR = 0.386 (this run; higher than the 0.174 reported in Exp 8 due to a different analogy-pair set and random seed).

k	Fixed (1500 ep)		Convergence		Matched loss	
	MRR	loss	MRR	epochs	MRR	epochs
2	0.248	0.050	0.113	491	0.120	283
4	0.158	0.050	0.392	608	0.304	176
8	0.195	0.050	0.375	526	0.254	176
16	0.286	0.050	0.290	607	0.306	178
32	0.421	0.050	0.330	498	0.168	172
64	0.246	0.050	0.300	727	0.156	188
128	0.311	0.050	0.215	431	0.269	185




Figure 5: Experiment 8b. *Left*: compositionality (analogy MRR) vs bottleneck dimension k under three stopping criteria. The $k = 2$ peak from the fixed-epoch condition disappears entirely under convergence and matched-loss controls; $k = 4$ dominates at convergence. *Right*: epochs to stopping criterion per k – smaller k converges in roughly half the steps of larger k , confirming the confound.

peaks instead), and it falls to the lowest MRR in both convergence and matched-loss conditions ($k = 2$ is last in both). Under convergence control, $k = 4$ shows the highest compositionality (0.392), above the raw LaBSE baseline (0.386); the matched-loss condition places $k = 16$ and $k = 4$ at the top.

The training-duration plot, confirms the confound: $k = 2$ converges in 283 epochs under matched-loss (the fastest by far), approximately half the ≈ 176 –608 epochs required by larger k . Within a fixed 1500-epoch budget, $k = 2$ receives roughly six times as many gradient steps relative to its effective convergence point as $k = 64$ does.

Two conclusions follow. First, the original $k = 2$ headline finding should be treated as fragile: it is sensitive to training duration and is not reproduced under either duration control. Second, the information-bottleneck interpretation remains plausible but requires revision: the compositionality sweet spot appears to be $k = 4$ rather than $k = 2$, and the effect is weaker than the original result suggested. A multi-seed replication with the matched-loss protocol is the appropriate next step.

Experiment 9: Universal Grammar Structural Probes

Motivation

Chomsky’s Universal Grammar (UG) claims that all human languages share deep structural properties emerging from a shared cognitive faculty. The hypothesis tested here is parallel but distinct: *do UG-like structural properties emerge from the geometry of the multilingual concept space without any explicit grammatical supervision?* If yes, a data-driven universal language would inherit these properties for free. If no, they must be explicitly designed in.

Setup

The experiment has three phases and five probes on 56 concepts across 5 languages (en, zh, es, ar, ru).

Phase 1 trains a lightweight projection layer on the LaBSE concept embeddings with a compositionality objective ($L_{\text{proximity}} + L_{\text{compositionality}}$), testing whether the space can be reshaped toward compositional structure.

Phase 2 builds a codebook of universal symbols (proto-tokens) from the projected centroids.

Phase 3 runs a Lewis signalling game in which agents must communicate concepts using 2-symbol messages drawn from the 56-symbol vocabulary, measuring whether communication pressure induces further structure.

The five probes test category emergence (Probe 1), argument structure (Probe 2), economy/Zipf’s law (Probe 3), cross-linguistic structure preservation (Probe 4), and semantic direction generalisation (Probe 5), in both the raw LaBSE space and the Phase 1 projected space.

Results

Phase 1: Compositional projection. The headline result of the entire experiment. Analogy MRR improves from 0.126 (raw LaBSE) to 0.451 after projection — a $3.6\times$ gain. Hits@1 goes from 0.0 to 0.333; Hits@3 from 0.333 to 0.667. This is not evidence that UG structure is spontaneously present in multilingual embeddings. It is evidence that compositional structure can be *induced* cheaply via gradient descent, reshaping the embedding space toward a geometry that supports semantic arithmetic. The Phase 1 training converges in under 800 epochs and the $L_{\text{compositionality}}$ loss approaches zero, while $L_{\text{proximity}}$ stabilises at 0.022, indicating that the projection stays close to the original concept positions while re-organising their relational structure.

Phase 2: Codebook. 56/56 entries used, 0 collisions. The projected space supports a clean, non-degenerate vocabulary of proto-tokens. New concepts (e.g. *ocean*, *hate*, *happiness*) map sensibly to nearest proto-tokens (*sky*, *love*, *joy*).

Phase 3: Lewis signalling game. Mean cosine reward converges to 0.745, below the 0.8 threshold. Crucially, symbol entropy *rises* monotonically over training rather than falling. Agents diversify their symbol usage rather than converging on a lean vocabulary. This replicates the anti-efficient encoding finding of Chaabouni et al. (2019): without a production cost (length or entropy penalty), agents exploit the full channel capacity. The concepts most resistant to communication are abstract social and aesthetic ones (memory, freedom, home, beauty, silence, reward 0.63–0.68), while emotional and action concepts communicate well (fear, hope, love, joy, war, reward 0.80–0.85). Adding a production cost penalty is the direct fix, and would connect Phase 3 to the opacity sweep of Experiment 3.

Table 10: Universal Grammar probe results. EMERGENT = property arises from information geometry alone; DETECTED = present after Phase 1 projection; NOT_DETECTED = requires explicit design.

Probe	Space	Metric	Verdict
P1: Category emergence	raw LaBSE	ARI=0.519	EMERGENT
P1: Category emergence	projected	ARI=0.433	EMERGENT
P2: Argument structure	both	LOO acc=0.625=chance	NOT_DETECTED
P3: Zipf/Economy	raw LaBSE	$R^2=0.708$	PARTIAL
P3: Zipf/Economy	projected	$R^2=0.855$	DETECTED
P4: Cross-linguistic RSA	universal	min $\rho=0.466$ vs max cross-lang=0.801	NOT_DETECTED*
P5: Semantic directions	raw LaBSE	opp acc=0.053	NOT_DETECTED
P5: Semantic directions	projected	opp acc=0.474	EMERGENT

Probe results.

Probe 1 (category emergence). K-means clustering ($k = 4$) recovers semantically coherent categories without labels (ARI=0.519): physical concepts (water, earth, wind, stone), action verbs (give, take, speak, think, run), abstract-emotional concepts (peace, memory, dream, home), and mixed clusters. This is robust across both spaces. The geometric structure of the concept space spontaneously encodes category information.

Probe 2 (argument structure). Leave-one-out logistic regression on the transitive/intransitive verb split achieves exactly 0.625 accuracy — equal to the class imbalance baseline (10 transitive / 6 intransitive). The geometric space encodes *no* transitivity signal. This is the clearest case for explicit design: if the universal language is to encode argument structure, it must be injected, not inherited.

Probe 3 (Zipf/economy). The projected space fits a power law with $R^2 = 0.855$ (vs 0.708 for raw LaBSE). The most central concepts in the projected space are fear, hope, war, love, peace — abstract but emotionally salient, which aligns with the expectation that high-frequency universals would be concepts with the most communicative pressure. The improvement from PARTIAL to DETECTED after projection confirms that the Phase 1 optimization is genuinely re-organising the geometric structure, not just shifting the space rigidly.

Probe 4 (cross-linguistic RSA). The NOT_DETECTED verdict requires careful interpretation. $D_{\text{universal}}$ has a minimum correlation with individual languages of $\rho = 0.466$ (vs Arabic), while some language pairs correlate at $\rho = 0.801$ (es–ru). The centroid space is not merely the average of the individual languages; it occupies a structurally distinct position. Whether this is a flaw (the centroid is a poor universal representation) or a feature (equidistance from all languages is exactly what universality requires) cannot be resolved from the current data. The key open question is whether $D_{\text{projected-universal}}$ — computed in the Phase 1 space rather than the raw centroid space — reverses the verdict. If the projection increases the minimum $D_{\text{universal}}-D_{\text{lang}}$ correlation above the maximum cross-language correlation, Probe 4 becomes EMERGENT.

Probe 5 (semantic directions). The opposition direction generalises to held-out antonym pairs only after Phase 1 projection (acc 0.053 $\rightarrow$ 0.474 on 19 held-out pairs). Correctly predicted pairs include light/dark, give/take, find/lose, cry/laugh, enemy/friend, pain/joy. Failing pairs include speak/silence and build/break, where the semantic opposition is more contextual than structural. This is the strongest argument for the Phase 1 projection as a grammar-induction

tool: the opposition structure that Esperanto and other constructed languages encode explicitly is learnable from analogy-pair supervision in the embedding space.



Figure 6: Top row: Phase 1 MRR improvement, Phase 3 Lewis game reward curve, Probe 1 category clusters. Bottom row: Probe 2 transitivity (NOT_DETECTED), Probe 4 RSA heatmap (NOT_DETECTED), Probe 5 semantic direction (EMERGENT in projected space). Bottom bar: overall UG scorecard, raw vs projected.

Analysis

The central result is that the universal language project cannot rely on raw multilingual embeddings alone for structure, but can induce substantial structure through cheap supervised projection. Phase 1 shows that compositionality, economy, and semantic directions are achievable with less than 800 gradient steps on a small set of analogy pairs. What is *not* achievable this way is argument structure (Probe 2): transitivity is a syntactic property that has no geometric footprint in concept embeddings trained on surface form. This is the property most in need of explicit grammatical design.

The Lewis game rising entropy result points to the most actionable next step: adding a production cost to Phase 3 should recover Zipfian symbol distribution (Galke & Raviv 2024), connecting the signalling game back to the opacity sweep of Experiment 3 and completing a circuit: discreteness (Exp 1) → optimal opacity (Exp 3) → production pressure (Exp 9 Phase 3) all converge on the same design principle for a universal communication system.




Figure 7: Probe 3: log-log fit of concept distances to power law. Left: raw LaBSE ($R^2 = 0.708$). Right: projected space ($R^2 = 0.855$). The projected space fits a steeper and cleaner power law, with emotionally salient concepts (fear, hope, war) as the most central.

Cross-Experiment Discussion

Several themes cut across all nine experiments.

Constraints improve symbol quality, but differently at different levels. Discreteness (Exp 1) enables error detection. Intermediate opacity (Exp 3) optimises task performance. Extreme compression (Exp 8) maximises compositionality. Production cost (absent in Exp 9 Phase 3, predicted to help) induces Zipfian structure. These are not the same mechanism, but they share a common principle: constraints select for the most structured solution consistent with task success, discarding arbitrary memorised associations. Deutsch (2011) expressed this as a general theory of knowledge growth through error correction; the experiments make it measurable.

The Phase 1 projection is the pivot of the universal language project. Without it, the concept space is rich but unstructured for compositional purposes (MRR=0.126, Probe 5 NOT_DETECTED). With it, the space supports semantic arithmetic (MRR=0.451), better Zipf structure (R^2 : 0.708→0.855), and generalising opposition directions (acc: 0.053→0.474). The Phase 1 optimization is cheap and does not destroy proximity to existing language words ($L_{\text{proximity}}$ stabilises at 0.022). This is the design tool for a data-driven universal language that was missing from v1.

English occupies an unusual position. In Experiments 6 and 7, English is consistently the most central language. In Probe 4 (RSA), English also has the highest correlation with $D_{\text{universal}}$ ($\rho = 0.818$). This is partly a consequence of LaBSE training distribution and speaker-count weighting, but it is also consistent with the claim in the PDF that English is the closest existing natural language to an interlingua.

Abstract social concepts are consistently the hardest. Experiment 4 finds the highest gap scores for aesthetic concepts. Experiment 9 Phase 3 finds the lowest communication rewards for memory, freedom, home, silence, beauty. These are not coincidentally the same concepts for which LaBSE centroids are most spread across languages. The implication for the universal language is that concrete nouns and primary emotions can be assigned tokens from centroid geometry alone; abstract social concepts require richer contextual definition.

Argument structure is the one property that cannot be induced. Probe 2 (transitivity) achieves exactly chance across both spaces. No other probe shows this degree of non-result. Transitivity is a structural property of how concepts participate in events, not a semantic property of the concepts themselves. The universal language will need explicit grammatical encoding

of argument structure, which is the task for Experiment 10.

Experiment 9b: Probe 4 in Projected Space and Production-Cost Lewis Game

Motivation

Two explicit open questions were left unresolved in Experiment 9.

Part 1. Probe 4 (cross-linguistic RSA) returned NOT_DETECTED with an asterisk: “The key open question is whether $D_{\text{projected-universal}}$ — computed in the Phase 1 space rather than the raw centroid space — reverses the verdict.” Experiment 9b computes Probe 4 in the projected space.

Part 2. Phase 3 of Experiment 9 found symbol entropy rising monotonically during Lewis-game training (anti-efficient encoding). The report identified adding a production-cost penalty as “the direct fix” to recover Zipfian symbol distribution and “complete the circuit: discreteness (Exp 1) $\rightarrow$ optimal opacity (Exp 3) $\rightarrow$ production pressure (Exp 9 Phase 3).” Experiment 9b tests whether a production cost achieves this.

Setup

Part 1. The same 56-concept, 5-language setup as Experiment 9. The Phase 1 compositional projection is re-trained (800 epochs), and Probe 4 cross-linguistic RSA is computed in both the raw centroid space and the projected space. The verdict criterion is: universal space is EMERGENT iff $\min_{\ell} \rho(D_{\text{universal}}, D_{\ell}) > \max_{\ell, \ell'} \rho(D_{\ell}, D_{\ell'})$.

Part 2. Two Lewis signalling games are trained for 600 epochs: one with no production cost, one with an entropy penalty on the symbol-usage distribution (weight 0.05). Final symbol-usage distributions are compared using Zipf power-law R^2 and symbol-usage entropy.

Results

Table 11: Probe 4 cross-linguistic RSA in raw vs projected space.

Space	$\min \rho(\text{univ}, \text{lang})$	$\max \rho(\text{cross-lang})$	Verdict
Raw centroid	0.686	0.831	NOT_DETECTED
Phase 1 projected	0.536	0.831	NOT_DETECTED

The gap between the minimum universal–language correlation and the maximum cross-language correlation is -0.145 in the raw space and -0.295 in the projected space. The projection *widens* the gap rather than closing it.

Table 12: Lewis game results with and without production cost.

Condition	Final entropy	Zipf R^2	Cosine reward
No production cost	3.350	0.817	0.723
With production cost	3.462	0.683	0.722

Analysis

Part 1 verdict: NOT FLIPPED. The Phase 1 projection does not reverse the Probe 4 verdict. Counterintuitively, the projection *reduces* the minimum universal–language correlation from 0.686 to 0.536 (making the universal space *less* similar to the most distant language). This can




Figure 8: Experiment 9b. *Top left*: Probe 4 bar chart — green bar (min univ-lang ρ) is below red bar (max cross-lang ρ) in both raw and projected spaces, confirming NOT_DETECTED. The projection narrows the green bar without closing the gap. *Top right*: Lewis-game symbol entropy over 600 epochs — both conditions show high-variance, high-entropy trajectories; no systematic entropy reduction from the production cost. *Bottom*: symbol-usage rank distributions — both are extremely concentrated (nearly all usage on symbol 0), giving high Zipf R^2 by default rather than from a true Zipfian mechanism.

be interpreted in two ways. Negatively: the Phase 1 projection optimises for compositionality, not equidistance from all languages, and these objectives are in tension. Positively: the Probe 4 NOT_DETECTED verdict may itself be a feature rather than a flaw — a genuinely universal space should be *equidistant* from all languages rather than highly correlated with any of them, and the max cross-language correlation of 0.831 simply reflects linguistic proximity between Spanish and Russian (or similar pairs) that a universal space should not reproduce.

The resolution of this ambiguity requires a richer criterion: rather than asking whether $D_{\text{universal}}$ correlates highly with individual languages, one should test whether it minimises the *variance* across per-language correlations. This is a direction for Experiment 10 or a revised Probe 4 formulation.

Part 2 verdict: NOT CONFIRMED. The production cost fails to recover Zipfian structure in this run. Symbol entropy is actually higher with the cost (3.462 vs 3.350) and Zipf R^2 is lower (0.683 vs 0.817). Inspection of the symbol distributions reveals a pathological regime: both games converge to near-degenerate solutions where one or two symbols receive almost all usage (symbols 0–2 dominate), making the rank distribution look Zipfian by construction — but the Zipf R^2 measure is dominated by these outlier counts rather than reflecting a genuine power-law vocabulary structure.

The production cost, applied as an entropy penalty on the usage distribution, *penalises* such collapse toward uniform usage, which paradoxically makes the distribution more uniform (higher

entropy) rather than more Zipfian. This is the wrong sign: a Zipfian distribution is non-uniform and skewed, so the entropy penalty actually pushes the game in the opposite direction. The correct intervention for recovering Zipf structure is a *length penalty* (penalising long messages, rewarding short ones) or a *frequency-based cost* (frequent symbols are cheaper), as in Galke & Raviv (2024). These are the next directions for Experiment 9b.

Experiment 10: Grammatical Frame Induction

Motivation

Experiment 9 Probe 2 found that transitivity (a binary verb property) is not recoverable from embedding geometry: leave-one-out logistic regression achieved exactly chance ($0.625 =$ class imbalance baseline). The report identified this as “the property most in need of explicit grammatical design.” However, transitivity is an unusually coarse probe — it collapses all argument-structure information into a single binary. Experiment 10 asks a finer question: do *frame roles* (AGENT, PATIENT, THEME, EXPERIENCER, INSTRUMENT, GOAL) have continuous geometric structure in the concept space, and does the Phase 1 compositional projection from Experiment 9 improve their separability?

If frame roles have geometric signal, they can serve as the empirical basis for agglutinative morphological suffixes in the universal language (analogous to Esperanto’s case system, but derived from data rather than a single designer’s intuitions).

Setup

56 concepts across 6 languages (en, es, de, fr, pt, ja) are annotated with FrameNet-inspired frame roles: AGENT (15), THEME (19), EXPERIENCER (12), STATE (5), INSTRUMENT (3), GOAL (2). Speaker-weighted LaBSE centroids are computed as in Experiment 9, and the Phase 1 compositional projection is re-trained (800 epochs).

Four probes are run in both raw LaBSE and projected spaces:

- **Probe A** (frame role classification): leave-one-out and 5-fold CV logistic regression on the 6-class role label. Chance baseline = 0.167.
- **Probe B** (continuous role geometry): binary LOO probes for `can_be_agent` (22/56), `can_be_patient` (22/56), `can_be_instrument` (8/56), plus Spearman ρ between each concept’s projection onto the learned role direction and its ground-truth binary label.
- **Probe C** (compositional frame test): given 10 verb+noun pairs, does the agent direction reliably identify which argument fills the agent role?
- **Probe D** (agent-patient asymmetry): do verb+agent compositions cluster separately from verb+patient compositions in PCA?

Results

Table 13: Probe A: frame role classification accuracy (chance = 0.167).

Space	LOO accuracy	5-fold CV accuracy
Raw LaBSE	0.554	0.539 ± 0.125
Projected	0.607	0.629 ± 0.162

Probe C. Raw LaBSE: agent identification accuracy = 0.900 (9/10). Projected: 1.000 (10/10). Both spaces support the compositional frame test; the projection closes the remaining error.

Table 14: Probe B: binary role geometry. Direction ρ = Spearman correlation between projection onto learned role direction and ground-truth binary label. All ρ values significant at $p < 0.0001$.

Role	Raw LaBSE		Projected	
	LOO acc	Direction ρ	LOO acc	Direction ρ
can_be_agent	0.661	0.821	0.750	0.819
can_be_patient	0.607	0.805	0.625	0.803
can_be_instrument	0.857	0.606	0.857	0.606

Probe D. Intra-agent coherence rises from 0.721 to 0.800 after projection; intra-patient coherence from 0.667 to 0.709. The asymmetry score (intra-agent coherence – inter-group similarity) increases from 0.024 to 0.030, confirming that verb+agent and verb+patient compositions form more distinct clusters in the projected space.

Analysis

Experiment 10 overturns the NOT_DETECTED verdict from Experiment 9 Probe 2. The key difference is resolution: Probe 2 tested a binary transitivity split that collapsed all role information; Probe A tests 6-class role prediction and finds above-chance accuracy in raw LaBSE (0.554) that improves further after projection (0.629). The direction ρ values from Probe B (0.80–0.82 for agent and patient) are particularly striking: the concept space contains a nearly monotone geometric signal for agent-hood and patient-hood even without any grammatical supervision.

Three design implications follow. First, agglutinative role suffixes for the universal language can be grounded in the agent and patient directions of the projected concept space — the suffixes would encode position along these learned geometric axes. Second, AGENT and THEME are cleanly separable (zero confusion in the projected-space confusion matrix); EXPERIENCER is the main source of error, which is linguistically expected (experiencer arguments share properties with both agents and patients). Third, the compositional frame test (Probe C) reaching 100% accuracy in the projected space means that, for verb+noun pairs, the universal grammar can assign roles from geometry alone — no separate syntactic parser is needed for the most common frame structures.

The one limitation is instrument: Probe B achieves only 0.857 for `can_be_instrument`, matching the majority-class baseline (8 positive / 48 negative). Instrument role is a genuinely sparse category, and more annotated examples would be needed before a data-driven instrument suffix could be proposed.

Experiment 12: TVS on Corpus Contexts (Tatoeba)

Motivation

Experiment 5 implemented the Translation Validity Score (TVS) on hand-crafted sentence pairs: each context was deliberately constructed to isolate a specific sense of a polysemous word, and the correct and wrong-sense translations were chosen by the experimenter. This controlled design validated the metric but left open whether it works on *naturally occurring* language. Experiment 12 replaces handcrafted contexts with sentences sampled from the Tatoeba parallel corpus, and replaces hand-chosen wrong-sense translations with generic machine translation (Google Translate via the `deep-translator` library) as the baseline. The key question is whether TVS, applied at corpus scale, still detects the gap between a sense-aware human reference translation and a context-agnostic MT system.

Experiment 10 — Grammatical Frame Induction
Do frame roles (agent/patient/theme/experiencer) have geometric structure?




Figure 9: Experiment 10. *Top row, left to right*: PCA of raw LaBSE centroids coloured by frame role; PCA of projected centroids (roles more separated); projection of each concept onto the learned agent direction (red = agent-capable, blue = patient-capable, grey = neither). *Bottom row*: binary role probe LOO accuracies (Probe B); frame role confusion matrix in the projected space (Probe A — AGENT and THEME nearly perfectly separated, EXPERIENCER the main source of confusion); PCA of verb+agent vs verb+patient compositions (Probe D — two clouds visibly separated along the first principal component).

Setup

Parallel sentence pairs are loaded from Tatoeba for **en→es** and **en→fr** via the HuggingFace datasets library. Where the dataset loader fails (as in this run, due to deprecated dataset scripts), a built-in fallback corpus of 44 sentence pairs (en-es) and 21 sentence pairs (en-fr) is used. Sentences containing the target polysemous word are filtered, sense clusters are inferred from source-sentence embeddings via k -means ($k = 3$), and diversity weights are assigned as distance from the cluster centroid (unusual contexts receive higher weight). MT translations are generated automatically; human references are drawn from the Tatoeba target-language column.

Six words are tested: *bank, light, run, play, set* (en→es) and *break, right, fall* (en→fr). A minimum of 5 usable context sentences per word is required to proceed; words below this threshold are skipped.

For each word, four quantities are computed per context: LaBSE cosine similarity between

source and human-reference target sentences (ref), LaBSE cosine similarity between source and MT target sentences (mt), diversity weight, and sense cluster membership. TVS is the diversity-weighted mean of the per-context similarities. The *weighting boost* is defined as (weighted gap) – (unweighted gap): a positive value means diversity weighting amplifies the human-over-MT advantage by penalising MT harder on rare-sense contexts.

Results

Four words reached the minimum-context threshold; three were skipped for insufficient corpus hits (*run*, *play*, *break*, *fall* each had 2–4 matching sentences).

Table 15: Experiment 12 TVS results on Tatoeba corpus contexts. Human ref TVS > MT TVS for 3 of 4 words. Weighting boost positive for 1 of 4 words.

Word	Lang pair	Contexts	Ref TVS	MT TVS	Weighted gap	Wt. boost
<i>bank</i>	en→es	12	0.900	0.900	+0.0007	−0.0010
<i>light</i>	en→es	10	0.898	0.900	−0.0023	−0.0001
<i>set</i>	en→es	5	0.879	0.872	+0.0072	−0.0006
<i>right</i>	en→fr	6	0.910	0.909	+0.0006	+0.0002
Mean across 4 words					+0.0015	−0.0004

Per-sense breakdown for *bank*. The most polysemous word in the set shows a clear sense-gap pattern. Three senses are detected: Sense 0 ($n = 4$), Sense 1 ($n = 7$, dominant), and Sense 2 ($n = 1$, minority). Human references maintain higher similarity than MT in Sense 0 (gap +0.007) and Sense 1 (gap +0.002), but MT outperforms the human reference on Sense 2 (gap −0.023). This is consistent with Sense 2 being a rare usage (aviation: “the plane banked sharply”) for which the Tatoeba human reference is suboptimal while generic MT, by coincidence, produces a closer paraphrase.




Figure 10: Per-sense translation quality for *bank* (en→es). Human reference maintains higher similarity on the two larger senses; MT scores higher on the single minority-sense context (Sense 2, $n = 1$), where the Tatoeba reference translation is non-optimal.

Analysis

Core result. TVS is positive (human > MT) for 3 of 4 words on naturally occurring corpus contexts, confirming that the metric generalises beyond hand-crafted sentences. The mean weighted gap across all four words is +0.0015 (human reference > MT), consistent in direction with Experiment 5.

Effect size. The gaps are substantially smaller than in Experiment 5 ($\Delta\text{TVS} = 0.119$ for *bank* in Exp 5 vs 0.0007 here). Two factors explain this. First, the hand-crafted contexts in Exp 5 were deliberately adversarial — each context forced a specific rare sense, maximising the




Figure 11: Per-context translation similarity for *bank* (sorted by diversity weight, most unusual contexts on the right). *Top*: human reference and MT similarity traces — the two are nearly overlapping across most contexts, with the human reference slightly higher on average. *Bottom*: diversity weights by sense cluster. The single Sense 2 context (pink, far right) receives the highest diversity weight (0.115), making it the most diagnostic context. Because MT outperforms the human reference on that context, the weighted gap is smaller than the unweighted gap — explaining the negative weighting boost for *bank*.

challenge for MT. Tatoeba sentences are not designed to stress-test polysemy; most occurrences of *bank* in casual parallel text involve the dominant financial sense, which MT handles well. Second, the Tatoeba human references are not expert translations; in rare-sense contexts they sometimes produce looser paraphrases than MT, compressing the gap or reversing it.

Diversity weighting. The weighting boost is positive for only 1 of 4 words (*right*, +0.0002) and negative for the other three. The mechanism is working correctly in the sense that minority-sense contexts receive higher weights — but the minority-sense MT translations are not consistently worse than the human references in naturally occurring text. The adversarial construction of Exp 5’s contexts was precisely what made the weighting effective there. Extending Exp 12 with additional, more deliberately chosen corpus sources (e.g. literary text or news where rare senses appear in sufficient quantity) would give the diversity weighting more signal to amplify.

Corpus size. The fallback corpus (44 en-es pairs) is very small; the most polysemous words (*run*, *play*) did not reach the 5-context minimum. A production run of this experiment requires the full Tatoeba dump (~10M sentence pairs) or a comparable large parallel corpus. At that scale, even rare senses of highly polysemous words would generate enough contexts for reliable per-sense TVS estimation.

Verdict. TVS scales to corpus contexts and continues to track human-over-MT sense coverage in the correct direction. The diversity-weighting mechanism requires naturally adversarial context distributions (rare senses present in sufficient numbers) to contribute meaningfully; with casual parallel text it remains neutral. The experiment validates TVS as a corpus-deployable metric and identifies minimum corpus size and sense diversity as the binding constraints for future large-scale runs.




Figure 12: Experiment 12 summary across all four words. *Left*: human reference and MT TVS side by side — gaps are small but consistently in the correct direction for 3 of 4 words. *Centre*: scatter of weighted vs unweighted gap — all four words cluster near the diagonal ($y=x$), indicating the diversity weighting has limited amplifying effect at this corpus size. *Right*: weighting boost per word — only *right* shows a positive boost; the others are slightly negative, driven by minority-sense contexts where MT happens to match better than the human reference.

Experiment 13: Blind Decoding Audit

Motivation

Experiment 8 reported that compressed vectors at $k = 16$ “decode cleanly to correct words across all four tested languages (en, ru, zh, ar) for all eight test concepts,” with cosine similarities of 0.96–1.00. This claim was tested on eight concepts *seen during encoder training*, which cannot distinguish genuine language-independent semantic content from memorised associations. The report called for a harder test: “generate compressed vectors for 50 held-out concepts, decode to nearest words, give only those word lists to a human (or model) and ask them to identify the concept. If the hit rate is high, the channel is genuinely interpretable.”

Experiment 13 implements this test with 20 held-out concepts — concepts *never seen during encoder training* — and tests identification across a sweep of bottleneck dimensions $k \in \{2, 4, 8, 16, 32, 64\}$.

Setup

The bottleneck autoencoder from Experiment 8 is trained on 40 seed concepts. 20 held-out concepts (song, dance, door, tree, mother, rain, anger, truth, journey, silence, wisdom, shadow, blood, bird, ocean, snow, sun, king, bread, moon) are compressed at each k and decoded to the top-3 nearest words in 4 languages (ar, en, es, zh) from an interpretability vocabulary of 109/79/80/60 entries respectively. An oracle identifier (the model itself, given only the decoded word lists and asked to guess the concept) is used as the auditor. Hit@1 = the concept is the oracle’s first guess; Hit@3 = the concept appears in the top-3 guesses.

Results

At $k = 16$, 0/20 concepts are correctly identified by the oracle at Hit@1 (the oracle’s first guess is wrong every time). Hit@3 = 0.250 means the correct concept appears in the top-3

Table 16: Blind decoding audit results across bottleneck dimensions.

k	Hit@1	Hit@3
2	0.000	0.050
4	0.000	0.050
8	0.050	0.150
16	0.100	0.250
32	0.100	0.250
64	0.150	0.250
chance (1/20)	0.050	0.150

for 5 of 20 cases. Representative failures reveal a systematic pattern: the decoder consistently retrieves high-salience training-set concepts (sky, earth, water, life, death) as nearest neighbours regardless of the actual held-out concept.




Figure 13: Experiment 13. *Left*: Hit@1 and Hit@3 identification rates vs bottleneck dimension k (dashed line = chance at 0.05 for Hit@1, 0.15 for Hit@3). Hit@1 is at or below chance for all $k \leq 32$. *Centre*: per-concept identification confidence at $k = 16$ (all red = all incorrect; confidence = cosine similarity of decoded words to oracle’s best-guess concept). *Right*: failure analysis — systematic confusion toward high-salience training concepts (sky appears 4 times as the wrong answer).

Analysis

The result is a clear revision of the interpretability claim from Experiment 8. Nearest-word decoding does *not* make the compressed channel human-auditable for held-out concepts. The fundamental problem is that the bottleneck encoder trained on 40 concepts maps the entire embedding space into a region spanned by those 40 training concepts. A held-out concept like *sun* is compressed to a region near the sky/light/earth cluster (the dominant training-set attractor), so its decoded word list reads “earth, sky, heaven” rather than “sun, moon, star.”

Two failure modes are distinguishable. The first is *training-set collapse*: concepts in the neighbourhood of several training concepts (e.g. *ocean* near water+earth) are decoded to the nearest training cluster regardless of their actual identity. The second is *compression aliasing*: at small k (2–4), multiple distinct concepts are mapped to the same bottleneck region and decoded identically.

The Hit@3 ceiling of 0.250 at $k \geq 16$ is meaningful: it means the correct concept is in the neighbourhood of the decoded word lists in about a quarter of cases. This is above chance (0.150) but far below the 0.70+ threshold that would justify the “human-auditable” claim. The Experiment 8 result (high cosine similarities for training concepts) was therefore an in-sample

artefact.

Two remedies are available. First, train the bottleneck encoder on a much larger concept set (500+) so that the training-set attractors cover more of the semantic space and held-out concepts land near correct neighbours. Second, replace nearest-word decoding with a dedicated decoder trained to produce concept labels from compressed vectors — this separates the compression quality from the interpretability interface.

Experiment 14: ZK Metric on Partially Communicating Systems

Motivation

Experiment 2 tested the ZK metric in a binary setting: a communicating game ($ZK = 0.851$) vs a severed baseline ($ZK = 0.000$). The report flagged that “the metric is most useful in the ambiguous middle” and called for a test on partially communicating systems to determine whether ZK tracks communication quality *continuously* or merely detects its presence. If continuous, ZK could serve as a training signal; if binary, it remains a diagnostic only.

Setup

Two sweeps are run, each training 60 epochs with the same two-turn referential game architecture as Experiment 2:

1. **Vocabulary size sweep:** $|V| \in \{2, 4, 8, 16, 32\}$ plus the severed baseline. Smaller vocabularies constrain the expressiveness of the channel and should reduce accuracy and, if ZK is continuous, ZK as well.
2. **Partial sender information sweep:** the sender observes $f \in \{2, 4, 8, 12, 16\}$ of the object’s 16 binary features. Fewer features = less information available to the sender = weaker communication signal.

After training, ZK and its three components (CIC_1 , RC, CIC_2) are computed. Spearman ρ between task accuracy and ZK score is the key diagnostic for continuity.

Results

Table 17: Vocabulary sweep results ($|V|$ = vocabulary size).

$ V $	Accuracy	ZK	CIC_1	RC	CIC_2
2	0.414	0.018	1.766	0.005	5.852
4	0.734	0.878	2.452	0.122	11.960
8	0.893	1.522	2.967	0.181	13.810
16	0.909	0.706	1.729	0.094	13.364
32	0.884	1.242	2.318	0.131	16.618
severed	0.203	0.000	0.072	0.000	0.492

Correlation analysis. Spearman $\rho(\text{accuracy}, ZK)$: vocab sweep = 0.400 ($p = 0.505$), partial sweep = 0.100 ($p = 0.873$), combined = 0.418 ($p = 0.201$). Neither sweep reaches statistical significance.

Analysis

Verdict: PARTIALLY CONTINUOUS. ZK correctly separates all communicating conditions from the severed baseline ($ZK = 0$ for severed in all runs), confirming its binary detection

Table 18: Partial information sweep results (f = features visible to sender).

$f/16$	Accuracy	ZK	CIC ₁	RC	CIC ₂
2/16	0.614	0.755	2.258	0.113	11.136
4/16	0.733	0.489	2.006	0.058	14.777
8/16	0.871	1.513	2.772	0.150	17.465
12/16	0.908	0.771	2.222	0.092	14.488
16/16	0.901	0.428	2.213	0.063	11.404

property from Experiment 2. However, the Spearman correlation with accuracy is positive but weak and non-significant in both sweeps ($\rho = 0.40$ and 0.10), so ZK cannot be used as a reliable continuous training signal.

The component analysis reveals *why* ZK is noisy as a continuous metric. RC (the sender-responds-to-query component) is the binding constraint on the ZK product and is nearly flat across all communicating conditions (range 0.058–0.181 across both sweeps). Small absolute changes in RC produce large proportional changes in the ZK product because ZK multiplies RC by the CIC average (which ranges from 7–10). A system that communicates well in both directions but happens to have a slightly less responsive sender (lower RC) will score lower than a system communicating less accurately but with a more responsive sender. This makes ZK sensitive to the sender–receiver balance rather than to overall communication quality.

The most actionable design revision is to replace the multiplicative ZK formula with an additive or min-pooled composite that does not amplify RC fluctuations:

$$\text{ZK}_{\text{revised}} = \alpha \cdot \text{CIC}_1 + \beta \cdot \text{RC} + \gamma \cdot \text{CIC}_2$$

with weights calibrated to equalise component variability (from the right panel of Figure 2). Given that CIC₂ is 8× more variable than CIC₁, setting $\gamma/\alpha = 1/8$ would balance their contributions. An alternative is a threshold-gated formula: $\text{ZK}_{\text{gated}} = \text{CIC}_2$ if $\text{RC} > \text{threshold}$, else 0, making RC a gating condition rather than a scaling factor.

The partial-information sweep confirms that ZK detects genuine communication even when the sender has limited information: $f = 2/16$ still produces $\text{ZK} = 0.755$, far above the severed baseline. This validates the metric’s robustness to degraded channel conditions.

sectionCross-Experiment Discussion

Thirteen experiments now span two distinct but connected research programmes: emergent communication under channel constraints (Exp 1–3, 1b, 3b, 9b, 14) and the structure of natural language as a window onto a universal concept space (Exp 4–10, 13). Several themes cut across both.

Constraints improve symbol quality, but the mechanism differs by level. Discreteness (Exp 1) enables error detection. Intermediate opacity (Exp 3, confirmed non-monotonic in Exp 3b) optimises task performance via accuracy-weighted RSA. Extreme compression (Exp 8) was initially claimed to maximise compositionality at $k = 2$, but Exp 8b showed this was a training-duration artefact; the true sweet spot is $k = 4$ under controlled convergence. Production cost in the Lewis game (Exp 9b Part 2) failed to recover Zipfian structure when implemented as an entropy penalty — the correct intervention is a length or frequency-based cost. The unifying principle remains: constraints select for the most structured solution consistent with task success. But the precise form of the constraint matters for which kind of structure is selected.




Figure 14: Experiment 14. *Top row*: vocabulary sweep — accuracy rises monotonically with $|V|$ and saturates at $|V| = 8$; ZK score is non-monotonic with a peak at $|V| = 8$, then drops at $|V| = 16$ and recovers at $|V| = 32$. Scatter plot (top right): ZK vs accuracy for the vocab sweep — the relationship is positive but with substantial noise ($\rho = 0.400$). *Bottom row*: partial information sweep — accuracy rises with features seen, ZK peaks at $f = 8$ and is non-monotonic. Combined scatter (bottom right): positive trend across both sweeps but p-value non-significant ($\rho = 0.418, p = 0.201$).

Frame roles have geometric structure; transitivity does not. Experiment 10 overturns the NOT_DETECTED verdict from Experiment 9 Probe 2 by using a finer probe. Direction ρ values of 0.82 for agent-hood and 0.80 for patient-hood are strong geometric signals, and Probe C achieves 100% agent identification in the projected space. The implication for the universal language is direct: role suffixes (agent, patient, theme) can be grounded in the geometry of the projected concept space, without explicit syntactic annotation. Transitivity as a binary verb property remains undetectable, but it can be inferred from role co-occurrence — a transitive verb is simply one that takes both an agent and a patient.

The interpretability claim from Experiment 8 requires revision. Experiment 13 found that the compressed channel is not human-auditable for held-out concepts: Hit@1 = 0.10 at $k = 16$, compared to the Exp 8 claim of near-perfect decoding. The gap is entirely explained by the in-sample/out-of-sample distinction: Exp 8 decoded concepts seen during training; Exp 13 decoded held-out concepts. The nearest-word decoder maps all held-out concepts toward the nearest training-set attractor. This does not invalidate the information-efficiency result (dense channel beats text at $k = 4$); it invalidates only the human-auditability result. The fix requires either a larger training set (500+ concepts) or a dedicated decoding head trained on concept labels.

The ZK metric is a reliable binary detector but a noisy continuous one. Experiment 14 confirms $ZK = 0$ for severed baselines across all conditions, validating the binary detection property established in Exp 2. However, Spearman $\rho(\text{accuracy}, ZK) = 0.42$ ($p = 0.20$) across both sweeps, which is insufficient for use as a training signal. The component analysis identifies RC (sender responsiveness) as the binding constraint and the source of noise: its low variability (std = 0.05 vs CIC_2 std = 3.14) makes the multiplicative ZK formula highly sensitive to small




Figure 15: ZK component analysis. *Left and centre*: CIC_1 , RC , and CIC_2 across vocab and partial sweeps. CIC_2 (receiver uses M_2) is the dominant and most variable component — it drives most of the ZK score variation. RC (sender responds to query) is the binding constraint and shows very low variability across conditions. *Right*: component variability (std across conditions) — CIC_2 is $\approx 8\times$ more variable than CIC_1 and $\approx 40\times$ more variable than RC , making it the primary diagnostic component.

RC fluctuations. A revised additive formulation with component-variance-equalising weights is the recommended next step.

English occupies an unusual position and the bias is structural. Across Exp 6, 7, 9, and 10, English is consistently the most central language — closest to the speaker-weighted centroid, highest Probe 4 RSA correlation, most natural interpretability-layer decoding target. Experiment 13’s failure mode is informative here: the training-set attractors that absorb held-out concepts are disproportionately English-labelled concepts (sky, earth, water, life, death). This suggests that LaBSE’s English-heavy training data creates attractors in the embedding space that are both linguistically and computationally dominant. Any universal language derived from this space will have an English bias that needs to be explicitly corrected for — either by downweighting English in the speaker-count weighting, or by regularising the projection to penalise over-concentration near English embeddings.

Abstract social concepts are the consistent failure mode. Exp 4 (highest gap scores for aesthetic concepts), Exp 9 Phase 3 (lowest Lewis-game rewards for memory, freedom, home, beauty), Exp 13 (held-out abstract concepts decoded to high-salience training attractors rather than accurate neighbours) all point to the same conclusion. Concrete nouns, primary emotions, and physical-world concepts can be handled by centroid geometry. Abstract social and aesthetic concepts require denser contextual annotation — more translation contexts (as in the TVS framework of Exp 5 and 12), more languages, or multimodal grounding.

The pipeline is now complete enough to identify the next binding constraint. The universal language project has: a draft vocabulary from centroid geometry (Exp 7, 11), translatability annotations per concept (Exp 4, 11), a translation quality metric (Exp 5, 12), geometric frame roles for morphology (Exp 10), a draft phonology (Exp 15), and ZK-verified communication metrics (Exp 2, 14). The binding constraint is no longer any single missing component but rather *scale*: every tool has been validated at 20–56 concepts and needs to be run at 500–5000 concepts to produce a usable language artifact. Experiment 11’s vocabulary compression pipeline is the natural entry point for this scale-up, with Experiment 10’s frame-role projections applied to the full compressed vocabulary.

Experiment 15: Phoneme Centroid Mapping

Motivation

The universal language pipeline has, at this point, a draft vocabulary grounded in multilingual embedding geometry (Exp 7, 11), translatability annotations (Exp 4), a translation quality metric (Exp 5, 12), and geometric frame roles for morphology (Exp 10). What is missing is *surface form*: the actual pronunciation of each proto-token. This is not cosmetic. The phoneme inventory constrains vocabulary size, learnability, and which existing languages the universal language will feel closest to for different speaker populations.

Experiment 15 addresses this by asking: if phoneme consensus is computed across the six most widely spoken languages (English, Spanish, French, German, Portuguese, Italian), weighted by speaker population, does the resulting proto-token inventory *naturally* satisfy the WALs near-universal phonological constraints — or must those constraints be imposed explicitly?

Setup

The 40-concept seed vocabulary from Experiment 11 is used. For each concept, the nearest word in each of six languages (en, es, fr, de, pt, it) is identified and converted to an IPA phoneme sequence via hand-crafted grapheme-to-phoneme rules (G2P), with an offline CMU-dictionary shim for English. The G2P pipeline produces 240 (concept, language) pairs with a mean of 5.1 phones per word and zero empty results.

Proto-token assignment proceeds as follows. For each concept, the union of phonemes across all six language words is collected and sorted by speaker-count-weighted frequency. The top four phonemes by weighted frequency are concatenated to form a four-phoneme proto-token. Coverage is defined as the fraction of the total speaker-weighted phoneme mass captured by those four phonemes; higher coverage means the proto-token will be phonetically familiar to more speakers.

Five WALs near-universal constraints are then checked against the global phoneme inventory: presence of stops (plosives), nasals, fricatives, vowels, and approximants. Finally, normalised phoneme edit distance between each proto-token and each language’s nearest word is computed for all 240 pairs, giving a per-language phonetic proximity ranking.

Results

Phoneme inventory. The global inventory contains 50 distinct segments across all proto-tokens, of which 35 belong to a named phonological class (the remainder are classified as “other,” principally the uvular trill R). The top segments by speaker-weighted frequency are the approximant /r/ (16.3), the vowels /a/ (14.3), /e/ (13.7), /i/ (11.8), and /o/ (11.5), and the nasal /n/ (10.7). By class, vowels dominate at 39.3% of the weighted inventory, followed by stops (20.4%), fricatives (12.0%), approximants (14.2%), nasals (9.7%), and other (4.3%).

WALS constraint check. All five WALs near-universal constraints are satisfied by the data-driven inventory without any explicit enforcement:

Proto-token draft vocabulary. Table 20 shows the top 15 concepts by coverage. The highest-coverage proto-token is TRUST (/tsna/, coverage 3.56), followed by DARK (/dkr/, 3.46) and BREAK (/rkbe/, 3.16). The five most phonetically natural concepts by mean normalised edit distance to existing words are HAND (0.583), NEW (0.592), LIGHT (0.610), TIME (0.611), and WIND (0.614). The mean edit distance across all 240 pairs is 0.764, indicating that the four-phoneme proto-tokens are on average substantially different from their source words —




Figure 16: Experiment 15. *Left*: global phoneme inventory of proto-tokens, coloured by phoneme class and sorted by speaker-weighted frequency. The approximant /r/ and the vowels /a/, /e/, /i/, /o/ dominate. *Right*: phoneme class distribution in the proto-token inventory — vowels account for 39.3%, stops for 20.4%, approximants for 14.2%, fricatives for 12.0%, and nasals for 9.7%.

Table 19: WALS near-universal phonological constraints in the Exp 15 proto-token inventory.

Constraint	Status	Representative segments
Stops (plosives)	PRESENT	/d/, /t/, /k/, /p/, /b/, /g/
Nasals	PRESENT	/n/, /m/, /ɲ/
Fricatives	PRESENT	/s/, /f/, /v/, /ʃ/, /h/, /ʒ/
Vowels	PRESENT	/a/, /e/, /i/, /o/, /u/, /ɜ/, ...
Approximants	PRESENT	/r/, /l/, /w/

expected given that a four-phoneme sequence drawn from the cross-language phoneme union will not in general reconstruct any single language’s word.

Per-language phonetic proximity. There is a dramatic asymmetry in how well the proto-tokens match existing language words. English has a mean normalised edit distance of 0.163 — the proto-tokens are substantially closer to English words than to any other language. German is the next closest (0.789), followed by Spanish (0.899), Italian (0.906), Portuguese (0.908), and French (0.921). The five Romance languages cluster tightly in the 0.899–0.921 range, while English is an outlier more than four times closer than the next language.

Analysis

Key finding: WALS universals emerge without explicit enforcement. All five WALS near-universal phonological constraints are satisfied by the speaker-count-weighted phoneme consensus. This is a direct answer to the experiment’s central question: a data-driven inventory derived from six major languages naturally contains stops, nasals, fricatives, vowels, and approximants, because these classes are present in every contributing language and therefore dominate the weighted phoneme pool. No explicit constraint-satisfaction step was needed.

The English proximity bias. The mean edit distance of 0.163 for English versus 0.789–0.921 for all other languages is the most striking quantitative result. It directly replicates and extends the finding from Experiments 6, 7, and 9 that English occupies an unusually central position in the multilingual embedding and phoneme spaces. The cause is likely twofold: (a) the English

Table 20: Top 15 proto-tokens by speaker-weighted phoneme coverage. Coverage = fraction of speaker-weighted phoneme mass captured by the four-phoneme proto-token. All 40 concepts processed across 6 languages.

Concept	Proto	Cov.	en	es	fr	de	pt	it
trust	/tsna/	3.56	trust	confianza	confiance	Vertrauen	confiança	fiducia
dark	/dkr/	3.46	dark	oscuridad	obscurité	Dunkel	escuridão	oscurità
break	/rkbe/	3.16	break	romper	casser	brechen	quebrar	rompere
river	/rivo/	3.14	river	río	rivière	Fluss	rio	fiume
time	/tma/	3.11	time	tiempo	temps	Zeit	tempo	tempo
hope	/paes/	3.06	hope	esperanza	espoir	Hoffnung	esperança	speranza
find	/nrfl/	3.01	find	encontrar	trouver	finden	encontrar	trovare
lose	/erlo/	2.86	lose	perder	perdre	verlieren	perder	perdere
change	/atnj/	2.78	change	cambiar	changer	ändern	mudar	cambiare
feel	/iefl/	2.71	feel	sentir	ressentir	fühlen	sentir	sentire
build	/irbl/	2.71	build	construir	construire	bauen	construir	costruire
sleep	/irls/	2.71	sleep	dormir	dormir	schlafen	dormir	dormire
friend	/afrn/	2.69	friend	amigo	ami	Freund	amigo	amico
light	/lta/	2.68	light	luz	lumière	Licht	luz	luce
wind	/ndw/	2.64	wind	viento	vent	Wind	vento	vento

G2P shim uses a CMU-dictionary lookup that produces compact IPA transcriptions matching the proto-token phoneme pool closely, while the rule-based G2P for other languages produces longer surface forms that diverge more; and (b) the speaker-count weighting assigns English a large share of phoneme mass (1.5 billion speakers), pulling the proto-token phonemes toward the English phoneme distribution. Any production deployment of this pipeline should explicitly downweight English in the speaker-count weighting — or use a larger, typologically balanced language sample — to avoid producing a proto-phonology that is effectively an English phoneme inventory.

The four-phoneme constraint. The mean edit distance of 0.764 across all language-concept pairs is high, reflecting the fundamental tension in condensing a 5–8 phoneme word into a 4-phoneme proto-token. The five most natural concepts (HAND, NEW, LIGHT, TIME, WIND) all have source words of 3–4 phonemes in at least one language, making the four-phoneme proto-token a closer approximation. Extending the proto-token length to five or six phonemes would reduce edit distances substantially for polysyllabic concepts; the optimal length is a hyperparameter that warrants a sweep.

Coverage as a design criterion. Coverage (the fraction of speaker-weighted phoneme mass captured by the proto-token) varies from 2.02 (FEAR) to 3.56 (TRUST). This variation is meaningful: high-coverage proto-tokens use the most frequent cross-linguistic phonemes and will therefore feel more familiar to more speakers. Low-coverage proto-tokens use rarer phonemes, potentially improving distinctiveness at the cost of learnability. The coverage metric is a natural objective for optimising the proto-token vocabulary once the pipeline scales to 500+ concepts.

Connection to the pipeline. Experiment 10 established that frame roles (agent, patient, theme) have geometric structure in the projected concept space and can serve as the empirical basis for agglutinative morphological suffixes. Experiment 15 provides the phoneme inventory from which those suffix phonemes should be drawn. Specifically, the most frequent segments in the proto-token inventory — /r/, /a/, /e/, /i/, /o/, /n/ — are the natural candidates for grammatical morphemes, as they appear across all six languages and have the highest speaker-weighted frequency. The least frequent segments (/ʃ/, /ʒ/, /ʒ/, /ʃ/) are better reserved for content distinctions, where their rarity helps maintain vocabulary distinctiveness.




Figure 17: Mean normalised phoneme edit distance between proto-tokens and nearest words per language. Lower = proto-tokens sound more like that language. English (0.163) is dramatically closer than all others; the five remaining languages cluster between 0.789 and 0.921.

Experiment A1: Master Vocabulary at Scale

Motivation

Every experiment in the universal language pipeline — Exp 4–10, 12, 13, 15, A2 — has operated on a vocabulary of 20–56 hand-picked seed concepts. The findings are theoretically grounded but practically limited: a 56-concept universal language covers roughly the same semantic range as a tourist phrasebook. Experiment A1 is the first scale-up: embed 5 000 frequency-list words across 12 languages, align them into cross-lingual concept clusters via LaBSE nearest-neighbour matching, compute speaker-weighted centroids, and annotate each concept with a translatability tier. The output files (`master_vocab.parquet`, `master_vocab_centroids.npy`) serve as the shared vocabulary for all downstream experiments (A2, B2, C1 and future work).

Setup

Languages and speaker weights. Twelve languages are included: English (en, 1.5B speakers), Chinese (zh, 1.1B), Hindi (hi, 0.6B), Spanish (es, 0.56B), Arabic (ar, 0.38B), Portuguese (pt, 0.26B), Russian (ru, 0.26B), French (fr, 0.28B), German (de, 0.13B), Japanese (ja, 0.125B), Korean (ko, 0.08B), Italian (it, 0.085B).

Vocabulary extraction. The top 5 000 lemmas by frequency are pulled from `wordfreq` for each language. After deduplication and removal of function words (closed-class items such as determiners, conjunctions, and pronouns), the English list retains 4 863 content words. Content-word counts vary by language due to script and tokenisation differences: Chinese yields only 251 unique lemmas at 5 000 words (many Chinese words are multi-character compounds), while Japanese (1 286) and Korean (761) similarly reflect agglutinative/logographic structure.

Alignment and centroid computation. Each English word is matched to its nearest neighbour in each other language via cosine similarity of LaBSE embeddings (the same pivot strategy as Exp 11). A speaker-weighted centroid is computed for each of the 4 863 English pivot words.

Concept clustering. UMAP (50 dimensions, cosine metric) reduces the 4 863 centroids, followed by HDBSCAN clustering. This collapses synonymy: words that align to the same region

of embedding space are grouped into a single concept cluster.

Gap score and tiers. The gap score (mean pairwise cosine distance across all 12 language translations) is computed for each concept. Four translatability tiers are assigned by gap score thresholds:

- **Universal** (< 0.05): 151 concepts
- **Near-universal** (0.05–0.20): 1 766 concepts
- **Culture-specific** (0.20–0.40): 2 601 concepts
- **Untranslatable** (> 0.40): 343 concepts

Results

Table 21: Experiment A1 summary statistics.

Metric	Value
Input words (English pivot)	4 863
Languages	12
HDBSCAN concept clusters	309
Unclustered (noise) concepts	852 (17.5%)
Mean gap score	0.2265
Universal (< 0.05)	151 (3.1%)
Near-universal (0.05–0.20)	1 766 (36.3%)
Culture-specific (0.20–0.40)	2 601 (53.5%)
Untranslatable (> 0.40)	343 (7.1%)

Most universal concepts. The ten lowest-gap concepts are: *new*, *love*, *you*, *all*, *one*, *news*, *http* (a function-word artefact), and the closed-class words *the*, *and*, *for*. Excluding function-word artefacts and the URL token, the most universal content words are *new* (gap = 0.005), *love* (0.007), *one* (0.010), and *news* (0.010). These are consistent with the Exp 6/7 finding that concrete, emotionally salient, and numerically grounded concepts cluster tightly across languages.

Hardest-to-translate concepts. The top 20 untranslatable concepts are almost entirely proper nouns and culturally specific abbreviations: *Hamilton*, *Kentucky*, *Cleveland*, *Clinton*, *Baltimore*, *Cambridge*, *Liverpool*, *Edinburgh*, *NFL*, *Minnesota*, *Detroit*, *Vegas*, *FBI*, *rugby*. This is methodologically expected: proper nouns and English-culture-specific institutions have no semantic neighbours in other languages’ frequency lists and thus produce maximally dispersed cross-lingual clusters.

Cluster compression. HDBSCAN identifies 309 concept clusters from 4 863 input words, a compression ratio of approximately $15.7\times$. This is substantially higher than the $2.0\times$ ratio in Exp 11’s 2 000-word sweep, reflecting that the 5 000-word list reaches deeper into the long tail where many near-synonyms cluster together (e.g. *begin*, *start*, *commence*, *initiate* collapse into one cluster). The 309-cluster figure is lower than the 2 000–4 000 expected range stated in the experimental design; this reflects a `min_cluster_size` parameter that is too large for the density of the long-tail region of the frequency distribution. Re-running HDBSCAN with a smaller `min_cluster_size` (e.g. 3 instead of the default) would recover more fine-grained clusters.

Analysis

The 151 universal concepts are the core vocabulary. These 151 concepts ($\text{gap} < 0.05$) can be assigned proto-tokens from centroid geometry alone, with high confidence that the result will be phonetically familiar and semantically transparent across all 12 languages. They include primary emotions, basic numerals, universal actions (eat, sleep, go, come), and physical-world primitives. Combined with the 1 766 near-universal concepts, the universal language has a working vocabulary of $\sim 1\,900$ concepts available without requiring culturally specific annotation.

Proper noun contamination requires filtering. The domination of the untranslatable tier by proper nouns (Hamilton, Kentucky, NFL, FBI, etc.) reveals that a naive frequency-list approach inherits English-centric named entities. A production vocabulary pipeline should apply a named-entity filter before computing gap scores. The proper noun cluster is identifiable post-hoc: it has very high gap scores (> 0.45) and most isolated language is consistently English (the only language where the word has a neighbour in the frequency list). Filtering these would reclassify ~ 200 of the 343 untranslatable concepts into the noise/excluded category and reduce the mean gap score.

The geographic structure of the translatability atlas is informative. The right panel of Figure 1 shows a clear spatial pattern: universal concepts cluster at the geometric centre of the UMAP, culture-specific concepts at the periphery, and untranslatable concepts as isolated outliers. This is the macro-scale version of the Exp 4 finding that gap score correlates with cultural specificity. It confirms that the translatability gradient is not random noise but a genuine geometric property of the multilingual embedding space: concepts that are easy to translate sit where all 12 language embeddings converge; concepts that are hard to translate sit where the language embeddings diverge.

Cluster count and HDBSCAN calibration. The gap between 309 observed clusters and the 2 000–4 000 expected is the main methodological finding to carry forward. The `min_cluster_size` hyperparameter controls the granularity of the synonymy collapse: too large, and semantically distinct concepts are merged; too small, and genuine synonyms are not collapsed. For a universal language vocabulary, the right criterion is not cluster count per se but concept distinctiveness: two concepts should be in the same cluster if and only if they have no context in which they are not interchangeable. A future version of A1 should calibrate `min_cluster_size` against a held-out synonymy judgement dataset.

Outputs and downstream use. The two output files produced by A1 are used throughout the remaining experiments:

- `master_vocab.parquet`: the full concept table with columns for English pivot word, cluster ID, gap score, translatability tier, and nearest-neighbour translations in all 12 languages.
- `master_vocab_centroids.npy`: the 768-dimensional speaker-weighted LaBSE centroids for all 4 863 concepts (14.9 MB). These centroids are the input to Experiment A2’s scaled UG probes, the Phase 1 compositional projection, and the frame-role geometry analysis.

The 151 universal and 1 766 near-universal concepts form the priority subset for Experiment A2 and for the first complete grammar draft in Experiment A3 (forthcoming).

Experiment A2: Scaled Universal Grammar Probes

Motivation

Experiments 9 and 10 established four probe results on 56 concepts across 5–6 languages: category emergence (EMERGENT, ARI=0.519), transitivity NOT_DETECTED, Zipf economy ($R^2=0.855$ in the projected space), and frame role geometry ($\rho = 0.82$ for agent-hood). All four were obtained at a scale that is arguably too small for reliable claims about the geometry of the concept space. Experiment A2 re-runs the same four probes at $10\times$ scale: 409 concepts (209 nouns, 200 verbs) across 5 languages (en, zh, es, ar, ru), with multilingual translations derived from LaBSE nearest-neighbour lookup rather than a multilingual WordNet. The central question is: do the Exp 9/10 findings hold?

Setup

A vocabulary of 500 candidate concepts is assembled from a curated noun and verb list, balanced across semantic categories (social, communication, motion, cognition, etc.). After filtering to concepts that embed cleanly in LaBSE and receive coherent nearest-neighbour translations in all five languages, 409 concepts remain (209 nouns, 200 verbs; 51 transitive verbs, 149 intransitive).

Speaker-weighted LaBSE centroids are computed as in Experiments 9–10. The Phase 1 compositional projection is trained for 600 epochs on 31 analogy pairs (derived from hand-coded antonym pairs and WordNet English antonyms). The same four probes are applied in both raw LaBSE and projected spaces.

Probe 1 (category emergence): k -means clustering, ARI against semantic category labels.

Probe 2 (transitivity): 5-fold CV logistic regression on the transitive/intransitive split across 200 verbs. Majority baseline = 0.745 (149/200 intransitive).

Probe 3 (Zipf/economy): log-log power-law fit of concept distances to centroid vs rank.

Probe 4 (frame role geometry): Spearman ρ between each concept’s projection onto the learned agent/patient direction and its ground-truth binary role label.

Results

Table 22: Experiment A2 probe results at 409-concept scale. Exp 9/10 values in parentheses for comparison.

Probe	Raw LaBSE	Projected	Exp 9/10 raw	Exp 9/10 proj
P1: ARI (category)	0.137	0.152	0.519	0.433
P2: CV acc (transitivity)	0.745	0.745	0.625	0.625
P3: Zipf R^2	0.861	0.915	0.708	0.855
P4: agent ρ	0.588	0.575	0.821	0.819
P4: patient ρ	0.600	0.589	0.805	0.803

Analysis

Probe 1 (category emergence): DEGRADED. ARI drops from 0.519 (56 concepts) to 0.137 (409 concepts). This is expected: k -means with a fixed k recovers coarser structure as the concept set grows, and the 20+ semantic categories in A2 are more fine-grained than the 4 broad clusters in Exp 9. It does not mean category structure disappears — the PCA plot shows clear

macro-regions — but it does mean that simple k -means ARI is not a stable metric across scales. A hierarchical clustering metric (e.g. dendrogram correlation) would be more appropriate at this scale.

Probe 2 (transitivity): CONFIRMED NOT_DETECTED. With 200 verbs and a 51/149 transitive/intransitive split, CV accuracy is exactly 0.745 in both raw and projected spaces — identical to the majority baseline. This is the most definitive result in the experiment: even at $4\times$ the Exp 9 verb count, embedding geometry carries no transitivity signal. The explicit grammatical design recommendation from Exp 9 is fully confirmed.

Probe 3 (Zipf/economy): STRENGTHENED. Both raw ($R^2 = 0.861$, up from 0.708) and projected ($R^2 = 0.915$, up from 0.855) Zipf fits improve at scale. The log-log plots show tighter fits with steeper slopes, and the most central concepts are high-frequency transitive verbs (make, give, take, find) rather than the emotionally salient concepts that dominated in Exp 9. This confirms that Zipfian economy is a robust property of the concept space that strengthens with vocabulary size.

Probe 4 (frame role geometry): PARTIALLY CONFIRMED. The agent and patient direction ρ values drop from 0.82/0.80 (Exp 10) to 0.59/0.60 here. Both remain highly significant ($p < 0.0001$) and substantially above zero, confirming that frame role geometry is present at scale. However, the decrease in ρ reflects two factors: (a) the larger, noisier annotation set (FrameNet labels applied automatically to 409 concepts rather than hand-annotated for 56); and (b) the broader vocabulary includes more concepts with ambiguous role assignments (e.g. stative verbs that are neither prototypically agentive nor patientive). The direction remains meaningful and the bar chart confirms that the top-30 concepts by agent projection are all unambiguously agent-capable verbs.

Summary verdict. The key Exp 9/10 findings scale qualitatively. Transitivity remains NOT_DETECTED. Zipf structure strengthens. Frame role geometry persists, though at lower ρ . Category emergence is scale-sensitive and requires a different metric above 200 concepts.

Experiment B2: Unified Constraint Grid Search

Motivation

Experiments 1–3, 8, and 9b each varied one constraint dimension in isolation: discreteness (temperature τ), compression (bottleneck k), and production cost. The cross-experiment discussion has repeatedly noted that these constraints likely interact — the optimal temperature at $k = 2$ may differ from the optimal temperature at $k = 32$, and adding a production cost may change the optimal message length. Experiment B2 runs a full factorial grid search over all four constraint dimensions simultaneously, asking: which constraint, or combination of constraints, produces jointly high accuracy *and* symbol quality (AW-RSA)?

Setup

A Lewis signalling game (binary object attributes, referential identification, 5-way choice) is trained for 80 epochs under 81 conditions: a $3 \times 3 \times 3 \times 3$ grid over:

- **Temperature** $\tau \in \{0.3, 1.0, 5.0\}$ (low = discrete, high = soft)
- **Bottleneck** $k \in \{2, 8, 32\}$ (compressed message dimension before discretisation)
- **Message length** $L \in \{1, 2, 4\}$ (number of symbols per message)
- **Production cost** $c \in \{0.0, 0.05, 0.2\}$ (entropy penalty on symbol usage)

For each condition, task accuracy and AW-RSA (accuracy $\times$ RSA) are measured. The Pareto frontier in the (accuracy, AW-RSA) plane identifies conditions that dominate — no other condition achieves higher values on both objectives simultaneously. Zipf R^2 is also recorded to test whether production cost recovers Zipfian structure, following the prediction from Exp 9b.

Results

Table 23: Top 10 conditions by AW-RSA across 81 grid conditions.

τ	k	L	cost	Accuracy	AW-RSA	Zipf R^2
1.0	32	4	0.20	0.959	0.315	0.419
1.0	32	4	0.00	0.945	0.314	0.698
1.0	32	4	0.05	0.944	0.288	0.494
5.0	32	4	0.20	0.892	0.285	0.796
1.0	8	4	0.20	0.920	0.283	0.775
5.0	32	4	0.05	0.896	0.276	0.791
0.3	32	4	0.05	0.920	0.271	0.913
1.0	8	4	0.00	0.920	0.269	0.701
5.0	32	4	0.00	0.905	0.257	0.969
5.0	8	4	0.20	0.878	0.250	0.924

Pareto frontier. Only 1 of 81 conditions lies on the Pareto frontier: $\tau = 1.0$, $k = 32$, $L = 4$, cost= 0.2 (accuracy=0.959, AW-RSA=0.315). This is the single condition that is not dominated — no other condition achieves higher accuracy *and* higher AW-RSA simultaneously.

Marginal effects. Mean AW-RSA by individual factor levels:

Table 24: Mean AW-RSA by factor level (marginalised over all other factors).

Factor	Level	Mean AW-RSA
Temperature τ	0.3	0.159
	1.0	0.175
	5.0	0.150
Bottleneck k	2	0.136
	8	0.169
	32	0.179
Message length L	1	0.082
	2	0.163
	4	0.238
Production cost c	0.00	0.159
	0.05	0.157
	0.20	0.167

Production cost and Zipf structure. Mean Zipf R^2 by production cost: 0.687 ($c = 0.0$), 0.700 ($c = 0.05$), 0.737 ($c = 0.2$). Higher cost produces higher Zipf R^2 , but the best single Zipf score is 0.969 at ($\tau = 5.0$, $k = 32$, $L = 4$, cost= 0.0) — a condition with no production cost. The Zipf structure in these cases arises from entropy collapse (the sender converges on a few dominant symbols), not from a genuine communicative economy pressure.

Analysis

Message length is the dominant constraint. The gap between $L = 4$ (mean AW-RSA = 0.238) and $L = 1$ (0.082) is 0.156 — larger than the gap from the best to worst level of any other factor (τ : 0.025; k : 0.043; cost: 0.011). Longer messages give agents more degrees of freedom to encode structural distinctions, and the combination with intermediate τ and large k provides the compositional capacity to use those degrees of freedom effectively. This result resolves the implicit question across Experiments 1–3, 8, and 9b: when all four constraint dimensions are varied jointly, message length wins.

Intermediate temperature remains optimal. $\tau = 1.0$ achieves the highest mean AW-RSA (0.175), consistent with the Exp 3b finding that an interior opacity level is optimal for accuracy-weighted symbol quality. $\tau = 0.3$ (near-discrete) and $\tau = 5.0$ (near-continuous) both underperform, confirming the non-monotonic opacity result across a much larger hyperparameter grid.

Larger bottleneck consistently helps. $k = 32$ outperforms $k = 8$ and $k = 2$ at every message length and temperature level. This is consistent with Exp 8b’s finding that $k = 4$ is the compositionality sweet spot under convergence control: within the range tested here ($k \in \{2, 8, 32\}$), more bottleneck capacity is always better because the 5-way task and 4-symbol message require more than 2 dimensions of representation.

Production cost has the smallest marginal effect on AW-RSA but the largest effect on Zipf structure. The cost factor spans only 0.011 AW-RSA units between its best (0.2) and worst (0.0) levels, confirming Exp 9b’s finding that an entropy penalty does not substantially improve symbol quality in the AW-RSA sense. However, mean Zipf R^2 rises by 7 points from cost= 0.0 to cost= 0.2. The increase in Zipf R^2 from production cost is genuine but modest, and is contaminated by the entropy-collapse regime at $\tau = 5.0$ where Zipf R^2 is high by default. The right interpretation is that production cost is a necessary but not sufficient condition for genuine Zipfian economy: it must be combined with a large bottleneck ($k = 32$) and long messages ($L = 4$) to produce a vocabulary structure that is both Zipfian and communicatively effective.

The optimal configuration for a universal communication system. The Pareto-efficient point ($\tau = 1.0$, $k = 32$, $L = 4$, cost= 0.2) provides an empirically grounded answer to the design question the report has been building toward. A practical universal communication protocol should use intermediate discretisation (not fully soft, not fully hard), a large enough bottleneck to support compositional encoding, messages long enough to express structural distinctions, and a mild production cost to encourage efficient vocabulary use. None of these settings is at an extreme, which is precisely what the non-monotonic opacity finding from Exp 3b predicted.

Experiment C1: Shannon Source/Channel Pareto Frontier

Motivation

The theoretical framing in the original PDF introduces Shannon’s source/channel coding separation theorem as the explanation for the tension between compression and error resilience: a communication system cannot simultaneously maximise both. Experiments 1–8 have repeatedly encountered this tension empirically — discrete protocols are compressible but fragile; continuous/dense protocols are resilient but bulky; text is somewhere in between. Experiment C1 operationalises Shannon’s theorem directly: it places all eleven communication protocols studied in this report on a two-dimensional compression–resilience plane and identifies the empirical Pareto frontier.

Setup

Eleven protocols are evaluated on a 30-concept referential identification task: raw LaBSE (float16, 1536 bytes), dense bottleneck encoders at $k \in \{2, 4, 8, 16, 32, 64\}$, Gumbel discrete at $\tau \in \{0.1, 1.0, 10.0\}$ (1 byte, vocabulary size 32), and text (natural-language description, mean 4 bytes per concept).

For each protocol, two quantities are computed:

- **Compression ratio:** $\log_2(N)/\text{message_bits}$ where $N = 30$ concepts. Higher = more compressed (more information per bit sent).
- **Noise resilience:** area under the accuracy-vs-noise curve across $\sigma \in \{0, 0.1, 0.2, \dots, 1.0\}$ (Gaussian noise for dense/text; token-flip noise for discrete). Higher = accuracy degrades more slowly.

The Pareto frontier in this plane identifies protocols that are not dominated: no other protocol achieves higher compression *and* higher resilience simultaneously.

Results

Table 25: All 11 protocols on the compression–resilience plane. Pareto-efficient protocols marked with $\star$.

Protocol	Msg bytes	Compression	Resilience	Pareto
Raw LaBSE (float16)	1536	0.000	0.549	
Dense $k = 64$ (128B)	128	0.005	0.614	
Dense $k = 32$ (64B)	64	0.010	0.614	
Dense $k = 16$ (32B)	32	0.019	0.614	
Dense $k = 8$ (16B)	16	0.038	0.614	
Dense $k = 4$ (8B)	8	0.077	0.611	
Dense $k = 2$ (4B)	4	0.153	0.617	
Text (4B avg)	4	0.153	0.828	$\star$
Gumbel $\tau = 0.1$ (1B)	1	0.613	0.082	
Gumbel $\tau = 1.0$ (1B)	1	0.613	0.077	
Gumbel $\tau = 10.0$ (1B)	1	0.613	0.086	$\star$

Two protocols are Pareto-efficient: **Text (4B avg)** and **Gumbel $\tau = 10.0$ (1B)**. Text achieves the highest noise resilience (0.828) at moderate compression (0.153). Gumbel $\tau = 10.0$ achieves the highest compression (0.613) at the cost of minimal resilience (0.086). No other protocol dominates either of these on both dimensions.

All dense protocols cluster tightly in a band (compression 0.005–0.153, resilience 0.611–0.617), with $k = 2$ matching text on compression but falling 13 points short on resilience. Raw LaBSE has the lowest compression of all (0.000) and moderate resilience (0.549).

Analysis

Shannon’s theorem is empirically visible. The Pareto frontier is not a point but a frontier: text achieves high resilience by sacrificing compression (using ~ 4 bytes for a natural-language description that a receiver can partially reconstruct even under noise); Gumbel achieves high compression by sacrificing resilience (1-byte discrete tokens that are destroyed by any noise). No protocol achieves both simultaneously. This is the empirical realisation of Shannon’s source/channel coding separation theorem for this referential task.

Dense protocols are not on the Pareto frontier. Despite the information-efficiency advantage established in Exp 8 ($k = 4$ beats text on bits-per-byte for the identification task), no

dense protocol appears on the Pareto frontier here. The reason is that the compression metric used in C1 ($\log_2 N / \text{message_bits}$) rewards protocols that use fewer bytes for the same number of referents, while the resilience metric rewards protocols that degrade gracefully under noise. Dense protocols at $k = 4$ (4 bytes) match text on compression but have resilience 0.611 vs text’s 0.828 — a gap of 0.217 that reflects the fundamental difference between a continuous embedding (sensitive to additive Gaussian noise) and a text description (partially reconstructible from partially corrupted tokens).

The dense protocol cluster’s near-uniform resilience across all k values (0.611–0.617) is itself a notable result: compression by a factor of 192 (from 1536 to 8 bytes) has almost no effect on noise resilience. All the resilience difference is between the channel type (continuous vs discrete vs natural-language) rather than the compression level within the continuous channel type.

Implications for the universal communication protocol. The Pareto analysis provides a precise answer to the question of which protocol type to use depending on the communication constraint. For noise-sensitive channels (adversarial, wireless, low SNR), text or a text-like protocol is optimal: its natural redundancy acts as an implicit error-correcting code. For bandwidth-constrained noiseless channels, discrete protocols with high compression are optimal. For the middle ground — moderate bandwidth, moderate noise — the dense protocols at $k = 4$ – $k = 8$ offer the best compression without the catastrophic fragility of fully discrete channels.

The connection to Experiments 1 and 3b is direct. Exp 1 showed that discrete channels fail detectably (rising wrong-answer entropy), while continuous channels fail silently. C1 quantifies the cost of that tradeoff across the full noise range: the area-under-curve resilience of 0.086 for Gumbel vs 0.828 for text represents an $\approx 10\times$ difference in noise tolerance. The question of which regime a universal language should target depends on its deployment context: human-to-human communication under cognitive noise (fatigue, distraction) favours the text-like high-resilience regime; AI-to-AI communication over a clean digital channel favours the dense high-compression regime.

Experiment A3: Multi-Hop Relay at 200 Concepts

Motivation

Experiment 8’s multi-hop relay used only 20 concepts and found near-perfect accuracy at every hop — a ceiling effect that made protocol comparison impossible. The report explicitly identified this as a gap: “follow-up experiments with 200+ concepts are needed to test whether dense or text degrades faster under multi-hop conditions with genuine information pressure.” Experiment A3 closes this gap by building a 700-concept vocabulary from LaBSE embeddings of frequent English words, training bottleneck autoencoders on 500 of them, and running a relay sweep on 200 held-out concepts over up to 8 hops at five noise levels.

Setup

A vocabulary of 700 English content words is embedded with LaBSE. Bottleneck autoencoders ($k \in \{4, 8, 16\}$, hidden 256, 2000 training epochs) are trained on the first 500 concepts with a combined cosine-reconstruction and referential identification loss. The held-out 200 concepts are used for all relay evaluations.

Five protocols are compared across $n_hops \in \{0, 1, \dots, 8\}$ and $\sigma \in \{0.0, 0.05, 0.1, 0.2, 0.3\}$:

- Dense $k \in \{4, 8, 16\}$: compress $\rightarrow$ add $\mathcal{N}(0, \sigma^2)$ to the k -dim bottleneck $\rightarrow$ decompress at each hop.

- Text (full emb): add Gaussian noise to the full 768-dim embedding at each hop (simulates token-level perturbation).
- Raw LaBSE: same as text but at half the noise rate (models the higher redundancy of the full embedding).

Identification accuracy = fraction of 200 held-out concepts correctly identified by nearest-cosine-neighbour lookup in the reference embedding space.

Results

Table 26: Identification accuracy at $\sigma = 0.1$, all hop counts.

Protocol	hop 0	1	2	3	4	5	6	7	8
Dense $k = 16$	1.000	0.575	0.440	0.285	0.130	0.070	0.035	0.025	0.005
Dense $k = 8$	1.000	0.450	0.345	0.205	0.155	0.100	0.060	0.040	0.035
Dense $k = 4$	1.000	0.225	0.145	0.085	0.095	0.045	0.030	0.030	0.020
Raw LaBSE	1.000	1.000	0.995	0.870	0.365	0.115	0.040	0.010	0.010
Text (full emb)	1.000	0.995	0.390	0.030	0.005	0.000	0.005	0.005	0.000

At $\sigma = 0.1$, hop 8: Dense $k = 8$ achieves accuracy 0.035, text 0.000, Raw LaBSE 0.010. The dense $k = 8$ advantage over text is +0.035, yielding the verdict: **similar overall degradation rate**, but with qualitatively different degradation curves.

Analysis

The ceiling effect from Exp 8 is resolved. At 200 concepts, every protocol degrades measurably within 3–4 hops at $\sigma \geq 0.05$, making protocol comparison meaningful.

Raw LaBSE is the most relay-resilient protocol by a large margin. It retains 99.5% accuracy at 2 hops and 87.0% at 3 hops under $\sigma = 0.1$. This is because the 768-dimensional embedding provides massive redundancy: noise applied to the full space has a small relative effect on the cosine similarity structure. The bottleneck compression that makes dense protocols information-efficient in the single-hop case (Exp 8) is precisely what makes them relay-fragile: noise in the k -dimensional bottleneck has a proportionally much larger effect.

Text collapses step-wise. The text protocol (simulated as additive Gaussian on the full embedding) drops from 0.995 at hop 1 to 0.390 at hop 2 and 0.030 at hop 3 under $\sigma = 0.1$. This matches the theoretical prediction: each re-encoding compounds the noise in a manner analogous to discrete re-quantization. The dense protocols degrade more smoothly, consistent with the prediction from the report, but the text collapse at hop 3 is qualitatively sharper than any dense protocol’s curve.

Compression bottleneck and text noise balance out at long chains. By hop 8, all protocols are effectively at noise-floor accuracy (0–4%). The dense advantage over text is small (+0.035 for best k) but positive. For a relay chain of practical length (3–6 hops), Raw LaBSE dominates despite its size disadvantage — a direct expression of the Shannon source/channel coding tension confirmed by Exp C1.

The connection to Exp C1 is direct: Raw LaBSE sits at the high-resilience end of the Pareto frontier (resilience 0.549 in C1) and shows the highest relay tolerance here. Dense protocols at $k = 4$ –16 sit in the intermediate band, while the text protocol’s steep step-wise collapse confirms its brittle reliance on the embedding structure surviving re-encoding.

Experiment B1: ZK Failure-Mode Taxonomy

Motivation

The ZK cross-experiment discussion noted that “a system with high CIC_2 but low RC is one-way signalling rather than genuine communication.” This claim implies that different communication failure modes should produce *distinct* (CIC_1 , RC, CIC_2) signatures — that the ZK metric can function as a diagnostic taxonomy rather than just a binary detector. Experiment B1 tests this by training five games, each with a different causal link deliberately broken, and verifying that each produces a recognisably different component pattern.

Setup

Five variants of the two-turn referential game from Experiment 2 are trained for 80 epochs (5-way choice, 16 binary features, vocabulary size 16, hidden 64):

- **Mode 0:** Full communication — all causal links intact.
- **Mode 1:** Receiver ignores M_1 — the query is zeroed before being sent to the sender; CIC_1 and RC should both approach zero.
- **Mode 2:** Sender ignores query — M_2 is a copy of M_1 regardless of the query; RC should approach zero.
- **Mode 3:** Receiver ignores M_2 — final prediction uses only M_1 ; CIC_2 should approach zero.
- **Mode 4:** Severed channel — both messages are zero vectors; the original Exp 2 severed baseline.

ZK components are computed on the validation set using 20 scrambling iterations per batch. Task accuracy is also recorded to confirm whether each mode can still solve the task through the remaining links.

Results

Table 27: ZK component signatures per failure mode.

Mode	Description	Acc	CIC_1	RC	CIC_2	ZK
0	Full communication	0.900	2.403	0.152	16.377	1.432
1	Receiver ignores M_1	0.910	0.464	0.000	18.707	0.002
2	Sender ignores query	0.764	0.373	0.000	5.079	0.000
3	Receiver ignores M_2	0.761	0.334	0.000	3.907	0.000
4	Severed channel	0.212	0.128	0.000	0.415	0.000

Analysis

The taxonomy works: each failure mode is distinguishable. Mode 0 (full communication) is the only condition with all three components elevated. Mode 1 (receiver ignores M_1) produces the specific pattern predicted by the report: $CIC_2 = 18.7$ (the highest of all modes) while $RC \approx 0$ and $CIC_1 \approx 0.46$. This is precisely one-way signalling — the receiver uses M_2 heavily, but the sender never responds to the query because the query carries no information about M_1 (which the receiver ignored). Modes 2–4 all converge to $RC \approx 0$, as predicted, but are distinguishable by their CIC_2 values: Mode 2 (sender ignores query) has moderate $CIC_2 = 5.1$ because the receiver still uses M_2 (which is just a copy of M_1); Mode 3 (receiver ignores M_2) has $CIC_2 = 3.9$ because a small residual effect of M_2 on prediction remains; Mode 4 (severed) has $CIC_2 = 0.4$, close to zero for both messages simultaneously.

Task accuracy is misleading without the taxonomy. Modes 0, 1, 2, and 3 all achieve 76–91% accuracy. A researcher evaluating only task performance would conclude all four games are communicating equally well. The ZK taxonomy reveals they are not: Mode 1 achieves 91% accuracy through one-way signalling alone (high CIC_2 , zero RC), while Mode 0 achieves 90% through genuine bidirectional communication.

The revised ZK formula recommended in Exp 14 is validated here. Under the multiplicative formula, Modes 1–4 all score near zero despite having non-trivial CIC values. The additive formulation $ZK_{\text{revised}} = \alpha \cdot CIC_1 + \beta \cdot RC + \gamma \cdot CIC_2$ would assign Mode 1 a substantially higher score than Mode 4, reflecting that it is closer to communication than a fully severed channel — which is the correct diagnostic conclusion.

Experiment B3: Phase 1 Projection Ablation

Motivation

Experiments 9 and 10 repeatedly invoke the Phase 1 compositional projection as a “cheap supervised projection.” The report never quantifies how cheap. Experiment B3 runs a systematic ablation: sweep the number of analogy pairs used for training ($n_pairs \in \{0, 2, 3, 5, 9, 15, 20\}$) and the $\lambda_{\text{compositionality}}$ weight ($\lambda \in \{0.25, 0.5, 1.0, 2.0, 4.0\}$) and measure analogy MRR and Probe 5 opposition accuracy at each combination.

Setup

The same 56-concept, 5-language vocabulary as Experiments 9–10 is used. 20 valid analogy pairs are available (drawn from antonym pairs where all four concepts appear in the vocabulary). For each (n_pairs, λ) combination, the projection is trained for 800 epochs. Raw LaBSE baseline: MRR = 0.334, Probe 5 = 0.250.

Results

Table 28: Analogy MRR as a function of analogy pairs and λ weight. Values uniform across λ unless noted.

n_pairs	$\lambda = 0.25$	$\lambda = 0.50$	$\lambda = 1.00$	$\lambda = 2.00$	$\lambda = 4.00$
0	0.334	0.334	0.334	0.334	0.334
2	0.632	0.632	0.632	0.632	0.629
3	0.675	0.675	0.675	0.675	0.676
5	0.825	0.825	0.825	0.825	0.825
9	0.904	0.904	0.904	0.904	0.904
15	1.000	1.000	1.000	1.000	1.000
20	1.000	1.000	1.000	1.000	1.000

The minimum number of pairs to exceed 50% of the maximum MRR gain is **3 pairs** (MRR = 0.675 vs raw 0.334; max 1.000). Probe 5 accuracy reaches 0.500 (above the baseline of 0.250) at $n_pairs = 3$ and is sustained at 9–15 pairs. Both metrics are completely insensitive to λ across the entire range tested.

Analysis

The “cheap supervision” claim is strongly confirmed — and cheaper than expected. Three analogy pairs (e.g. *love/fear*, *life/death*, *give/take*) are sufficient to achieve more than half the maximum MRR gain. Nine pairs achieve MRR = 0.904, accounting for 85% of the gap to saturation. The projection used in Experiments 9 and 10 (8–9 pairs by design) falls near

the knee of the saturation curve — a well-chosen setting, though in retrospect even fewer pairs would have sufficed for the headline results.

λ is irrelevant within the tested range. MRR is identical across all five λ values for every $n_{\text{pairs}} \leq 9$, and differs by at most 0.001 at larger counts. This means the proximity loss term (which keeps the projection close to the original centroid positions) completely dominates the geometry at low pair counts, and the compositionality loss only “fills in” structure once there are enough pairs to constrain the direction of improvement. For practical use, $\lambda = 1.0$ (the Exp 9 default) is as good as any other value in the range.

The practical recipe: 5 antonym pairs + $\lambda = 1.0$ + 800 epochs achieves $\text{MRR} = 0.825$ on the 56-concept vocabulary. Scaling to the A1 master vocabulary (4863 concepts) with proportionally more pairs (say, 50–100 antonym pairs from WordNet English) should produce comparable or better results, since the larger vocabulary provides a more discriminative test and the Phase 1 projection has been shown (Exp A2) to remain effective at 409 concepts.

Experiment C2: Language-Design Loss

Motivation

The theoretical PDF directly specifies a multi-term loss for constructing the universal language: “a measure of dissimilarity between the new language and the corpus of human language... so that it’s 2 times closer to English than to Spanish if there is 2 times more written English... [plus] the ability of the new language to express the interlingua space of ideas, without compromising on complexity.” Experiment C2 operationalises this three-term loss, optimises a learnable proto-vocabulary against it, and shows the coverage–simplicity trade-off curve. This is the closest the project comes to *constructing* the language rather than measuring properties of existing languages.

Setup

A concept space of 56 concepts is embedded with LaBSE across 10 languages (en, zh, hi, es, ar, ru, pt, fr, de, ja), weighted by speaker counts. A speaker-weighted centroid matrix (56×768) is computed.

A proto-vocabulary of N learnable 768-dimensional unit vectors is optimised against:

$$\mathcal{L} = \lambda_{\text{prox}} \cdot L_{\text{prox}} + \lambda_{\text{expr}} \cdot L_{\text{expr}} + \lambda_{\text{simp}} \cdot L_{\text{simp}}$$

where:

- L_{prox} (proximity): $1 -$ mean cosine similarity between each proto-token and its nearest concept centroid. Encourages proto-tokens to sit in the existing concept space.
- L_{expr} (expressivity): fraction of concepts with no proto-token within cosine similarity 0.3 (coverage penalty).
- L_{simp} (simplicity): N/N_{concepts} , penalising large vocabularies.

The sweep covers $N \in \{10, 20, 50, 100, 200\}$ with weights $\lambda_{\text{prox}} = 1.0$, $\lambda_{\text{expr}} = 2.0$, $\lambda_{\text{simp}} = 0.5$, trained for 500 epochs each.

Results

All vocabulary sizes achieve full coverage (1.000) after optimisation — the entire concept space is covered even by 10 proto-tokens. Both L_{prox} and L_{expr} reach zero for all N ; the total loss is driven entirely by L_{simp} , which grows linearly with N .

Table 29: Language-design loss results across vocabulary sizes.

N	L_{total}	L_{prox}	L_{expr}	L_{simp}	Coverage
10	0.088	0.000	0.000	0.175	1.000
20	0.175	0.000	0.000	0.351	1.000
50	0.439	0.000	0.000	0.877	1.000
100	0.877	0.000	0.000	1.754	1.000
200	1.754	0.000	0.000	3.509	1.000

Analysis

The optimisation succeeds but reveals an important structural limitation. L_{prox} and L_{expr} both reach zero for all N , meaning the gradient descent correctly finds proto-vocabularies that are close to the centroid space and cover all concepts. The Pareto-optimal vocabulary size in this experiment is $N = 10$, which achieves full coverage at the lowest total loss.

The result is dominated by the concept-space geometry, not the vocabulary. Full coverage at $N = 10$ for 56 concepts means the LaBSE centroid space is clustered enough that 10 prototype vectors suffice to be within cosine 0.3 of every concept. This reflects the well-known geometric property of LaBSE embeddings — concepts are not uniformly distributed in 768-dimensional space but cluster tightly around a few semantic macro-regions. The practical implication is that the coverage threshold (0.3 here) is the controlling parameter: tightening it to 0.5 or 0.7 would reveal a more meaningful vocabulary-size trade-off.

The simplicity term is the correct binding constraint. The fact that L_{simp} is the only non-zero term confirms the PDF’s intuition that the primary cost of a universal language is its description length, not its coverage. Any vocabulary large enough to span the concept space will achieve full expressivity; the interesting design question is how small the vocabulary can be made while maintaining adequate coverage at a reasonable coverage threshold. The next step is to re-run the sweep with adaptive coverage thresholds (e.g. cosine ≥ 0.6 , ≥ 0.7 , ≥ 0.8) to find the actual knee of the coverage–simplicity trade-off.

Experiment C3: MI Entanglement Metric

Motivation

Section 8 of the theoretical PDF states: “communication is entanglement. . . mutual information of each agent’s world models increases [as they develop a shared language]. The layer of theory of mind between these agents makes it so that socializing is a shared journey through the world.” This claim has remained entirely verbal — no experiment in the project has measured the mutual information between agent world models during training. Experiment C3 operationalises the claim: train two agents with *separate* world-model encoders (no parameter sharing), measure the InfoNCE lower bound on their mutual information over training, and test whether MI rises as a shared communication protocol emerges from task pressure alone.

Setup

Two agents are trained on a 10-way referential game (16 binary features). Each agent has its own world-model encoder (16→64→32 WORLD_DIM) and a message head (32→16 MSG_DIM). At each step, Agent 1 encodes the object and sends a message; Agent 2 receives it and must identify the target (and vice versa). There is no direct loss on the world-model representations — only the referential task loss.

Mutual information is estimated using InfoNCE: a learned projector maps each agent’s world-model state z_i to a normalised projection, and InfoNCE provides the lower bound $\text{MI}(z_1, z_2) \geq \log B - \mathcal{L}_{NCE}$ where B is the batch size. A severed baseline trains a single agent on self-communication (same architecture, no cross-agent interaction) and computes MI against a random vector, establishing the floor.

Results

MI is measured every 5 epochs over 100 training epochs ($N = 8000$ training objects, batch size 128).

Table 30: MI and accuracy at representative epochs.

Epoch	Task accuracy	MI lower bound	Severed MI
5	0.986	−0.310	−1.2 (approx.)
25	0.995	−0.294	−1.2 (approx.)
50	0.995	−0.278	−1.2 (approx.)
75	0.996	−0.270	−1.2 (approx.)
100	0.996	−0.265	−1.2 (approx.)

MI rises monotonically from −0.310 at epoch 5 to −0.265 at epoch 100, a net rise of +0.045. The severed baseline MI fluctuates around −1.2 to −1.4 throughout training with no upward trend (net change −0.166). Task accuracy converges rapidly (99.5% by epoch 20) and stabilises; MI continues rising slowly after accuracy plateaus, confirming that MI improvement is not simply a proxy for better task performance.

Analysis

Verdict: MI RISES with communication — entanglement is real. The MI lower bound rises by +0.045 nats for communicating agents vs −0.166 for the severed baseline (a relative difference of +0.21 nats). This directly confirms the PDF’s prediction: developing a shared communication protocol increases the mutual information between the agents’ separate world models, even without any direct MI optimisation objective and without parameter sharing.

The effect size is modest. A rise from −0.310 to −0.265 is numerically small, and the absolute MI values remain negative (the InfoNCE estimator reports a lower bound, and negative values simply indicate the bound is not tight). The communicating agents’ MI is consistently ~ 0.9 nats above the severed baseline at all epochs, which is the meaningful comparison. The continuing slow rise after epoch 50 (when accuracy is already at 99.5%) suggests that the shared language continues to refine structurally even after task performance saturates — consistent with the “shared journey through the world” claim in the PDF.

The MI trajectory reveals a two-phase structure. Task accuracy rises rapidly (epoch 5–20), then plateaus. MI rises more slowly throughout, with the steepest increase in the 25–50 epoch range after accuracy has stabilised. This is consistent with the interpretation that the agents first establish a functional protocol (accuracy phase) and then gradually align their internal representations more precisely (MI phase) — a form of implicit world-model synchronisation driven purely by communication pressure.

Connection to the ZK taxonomy (Exp B1) and entanglement. The ZK metric measures whether communication is causally bidirectional; C3 measures whether it is informationally integrative. Together they characterise two orthogonal properties of “genuine” communication: causal (both agents influence each other’s behaviour) and informational (both agents come

to represent the world more similarly). A complete characterisation of AI-AI communication quality requires both metrics.

Cross-Experiment Discussion

Twenty-three experiments now span two distinct but connected research programmes: emergent communication under channel constraints (Exp 1–3, 1b, 3b, 9b, 14, B1, B2, B3, C1, C3) and the structure of natural language as a window onto a universal concept space (Exp 4–15, A1, A2, A3, C2). Several themes cut across both.

Constraints improve symbol quality, but message length matters most. Experiment B2’s full factorial grid search is the definitive answer: when τ , k , L , and production cost are all varied simultaneously, *message length* produces the largest marginal effect on AW-RSA (gap 0.156 between $L = 4$ and $L = 1$, vs 0.043 for k and 0.025 for τ). The optimal configuration ($\tau = 1.0$, $k = 32$, $L = 4$, cost = 0.2) is the single Pareto-efficient point across 81 conditions. No setting is at an extreme, which is precisely what the non-monotonic opacity finding from Exp 3b predicted.

Shannon’s theorem is the unifying theoretical limit. Experiment C1 confirms that no protocol dominates on both compression and resilience simultaneously. Experiment A3 extends this to multi-hop relay: Raw LaBSE has the highest relay resilience but the worst compression; dense protocols at $k = 4$ –16 sit in an intermediate band; text collapses step-wise after 2–3 hops. The relay behaviour at 200 concepts is qualitatively consistent with the Pareto structure: high-redundancy protocols (Raw LaBSE) degrade slowly; compressed protocols degrade fast.

The Phase 1 projection is cheap — exactly 3–5 pairs cheap. Experiment B3 quantifies what “cheap supervision” means: 3 antonym pairs achieve 50% of the maximum MRR gain; 9 pairs achieve 85%; 15 pairs saturate at $\text{MRR} = 1.0$ on the 56-concept vocabulary. The λ weight is irrelevant across the entire tested range. For the A1-scale vocabulary (4863 concepts), proportionally scaling to 50–100 pairs should replicate the qualitative finding. This makes the Phase 1 projection the most practically deployable tool in the pipeline: a 5-minute annotation effort produces the full grammatical structure induction benefit.

The ZK metric is a diagnostic taxonomy, not just a binary detector. Experiment B1 confirms the report’s claim that different failure modes produce distinct (CIC_1 , RC, CIC_2) signatures. The one-way signalling signature (high CIC_2 , near-zero RC) correctly identifies Mode 1 (receiver ignores M_1) as distinct from full communication despite identical task accuracy (91% vs 90%). This validates ZK as a *diagnostic instrument*: a researcher can infer the structural failure mode from the component pattern without ever decoding message content. The revised additive ZK formula is further supported by B1: under it, Mode 1 would score substantially above the severed baseline, correctly reflecting that one-way signalling is closer to communication than complete severance.

Communication creates entanglement: MI rises with shared language. Experiment C3 directly confirms the PDF’s prediction: MI between two agents’ separate world models rises monotonically during training (+0.045 nats vs -0.166 for the severed baseline), and continues rising after task accuracy saturates. The two-phase structure (rapid accuracy convergence, slow MI refinement) suggests agents first establish a functional protocol and then gradually synchronise their internal representations. This is the “shared journey through the world” made measurable.

The language-design loss bottoms out at coverage, not expressivity. Experiment C2 finds that even 10 proto-tokens achieve full concept coverage on the 56-concept space, because LaBSE centroids cluster tightly enough that a small set of prototypes spans the space within